

# Credit Card Fraud Detection: Using a Generative Adversarial Network (GAN) to Generate Synthetic Data to Resolve Class Imbalance

Yu chien (Calvin) Ma

- In certain datasets, like credit card fraud detection, one class (e.g., non-fraudulent transactions) often greatly outnumbers the other (fraudulent transactions), leading to class imbalance. This imbalance can introduce biases when training machine learning models, as the model may become overly biased toward predicting the majority class.
- In this project, I first applied Principal Component Analysis (PCA) to reduce the dimensionality of the dataset, which originally contained 29 features, down to 2 components. This dimensionality reduction helps to visualize the structure of the data, revealing how fraudulent transactions compare to genuine ones in a lower-dimensional space.
- Next, I implemented a Generative Adversarial Network (GAN) to generate synthetic fraudulent transactions. The GAN consists of two components: a Generator that creates new, synthetic fraudulent data, and a Discriminator that evaluates the authenticity of the generated data against real fraudulent transactions. By training these two components in opposition, the GAN is able to generate increasingly realistic fraudulent transactions.
- Finally, I visualized the comparison between the real and synthetic fraudulent transactions across the 28 features. This allowed me to analyze how each feature behaves in both the genuine and synthetic fraud samples, providing insights into the characteristics of fraudulent transactions and how well the GAN can replicate those characteristics.

## Importing the Dataset

```
In [80]: import numpy as np
import pandas as pd

# Importing neural network modules
import tensorflow as tf
from tensorflow.keras.layers import Input, Dense, BatchNormalization, LeakyF
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.initializers import RandomNormal
# Importing some machine learning modules
```

```

from sklearn.utils import shuffle
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
# Import data visualization modules
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px

```

Check the data shape

```

In [82]: data = pd.read_csv("Creditcard_dataset.csv")
data.head()

```

```

Out[82]:

```

	Time	V1	V2	V3	V4	V5	V6	V7
0	1	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
1	7	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631
2	10	1.449044	-1.176339	0.913860	-1.375667	-1.971383	-0.629152	-1.423236
3	10	0.384978	0.616109	-0.874300	-0.094019	2.924584	3.317027	0.470455
4	11	1.069374	0.287722	0.828613	2.712520	-0.178398	0.337544	-0.096717

5 rows × 31 columns

Display number of genuine and fraudulent records we have (0 means genuine, and 1 means fraudulent)

```

In [8]: data.Class.value_counts()

```

```

Out[8]:
Class
0    50000
1      492
Name: count, dtype: int64

```

## Data Preprocessing and Exploration

- Removing all the rows with **Nan** values
- Removing **Time** column
- Feature Scaling **Amount** column
- Split the data into features and labels
- Data Exploration

Removing the rows **Nan** values in the dataset

```

In [11]: data = data.dropna()

```

Removing Time column

```
In [13]: data = data.drop(axis=1, columns = "Time")
```

Feature Scaling of Amount column

```
In [15]: data
```

```
Out[15]:
```

	V1	V2	V3	V4	V5	V6	V7
0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
1	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631
2	1.449044	-1.176339	0.913860	-1.375667	-1.971383	-0.629152	-1.423236
3	0.384978	0.616109	-0.874300	-0.094019	2.924584	3.317027	0.470455
4	1.069374	0.287722	0.828613	2.712520	-0.178398	0.337544	-0.096717
...	...	...	...	...	...	...	...
50487	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850
50488	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170
50489	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739
50490	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002
50491	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050

50492 rows × 30 columns

```
In [16]: # to scale the "Amount" column to match the other features as large, unscale  
  
scaler = StandardScaler()  
data["Amount"] = scaler.fit_transform(data[["Amount"]])
```

```
In [17]: data.head()
```

```
Out[17]:
```

	V1	V2	V3	V4	V5	V6	V7	
0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.37
1	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631	-3.80
2	1.449044	-1.176339	0.913860	-1.375667	-1.971383	-0.629152	-1.423236	0.04
3	0.384978	0.616109	-0.874300	-0.094019	2.924584	3.317027	0.470455	0.53
4	1.069374	0.287722	0.828613	2.712520	-0.178398	0.337544	-0.096717	0.11

5 rows × 30 columns

Splitting the genuine and fraud records into separate dataframes

```
In [19]: data_fraudulent = data[data.Class == 1]
data_genuine = data[data.Class == 0]
```

```
In [20]: data_fraudulent.shape
```

```
Out[20]: (492, 30)
```

```
In [21]: data_genuine.shape
```

```
Out[21]: (50000, 30)
```

Split the data into features and labels

```
In [23]: X = data.drop(axis = 1, columns = "Class")
Y = data.Class
```

Data Exploration

- Apply PCA to reduce the dimensionality of features `X` into two dimensions
- Reduce from 29 to 2 dimensions
- Use a scatter plot to visualize our data

```
In [25]: pca = PCA(2)
df = pd.DataFrame(pca.fit_transform(X))
```

```
In [26]: df.head()
```

```
Out[26]:
```

	0	1
0	0.447840	-1.197485
1	0.582393	-0.258062
2	0.939390	0.728299
3	0.630766	0.499103
4	0.536287	1.055403

```
In [27]: df["label"] = Y
```

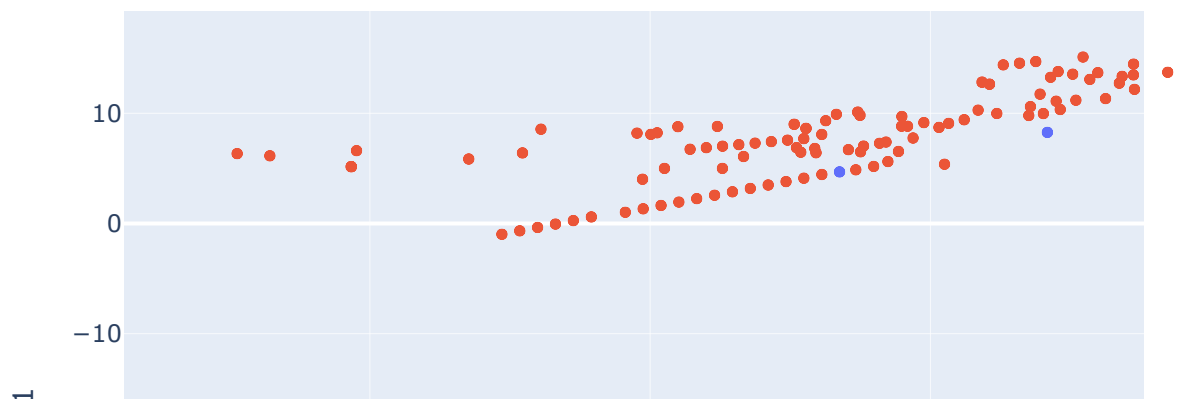
```
In [28]: df.head()
```

Out [28]:

	0	1	label
0	0.447840	-1.197485	0
1	0.582393	-0.258062	0
2	0.939390	0.728299	0
3	0.630766	0.499103	0
4	0.536287	1.055403	0

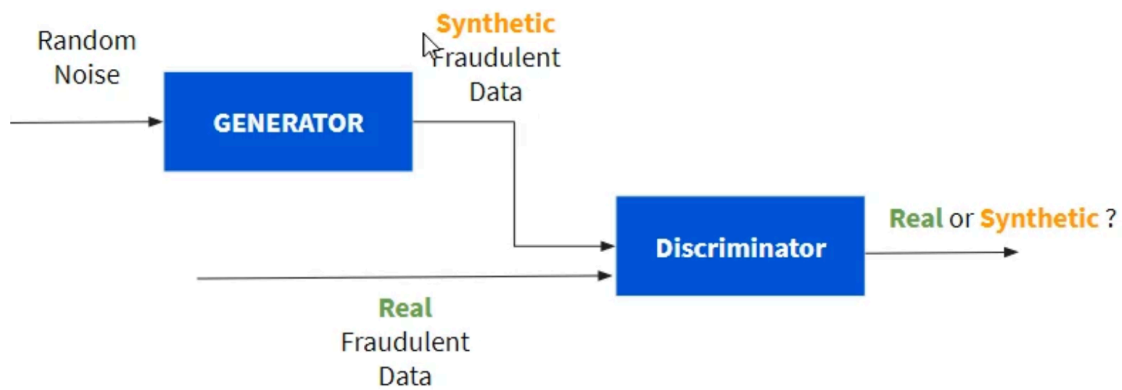
Plotting the two PCA features to visualize the genuine and fraudulent transactions

```
In [30]: px.scatter(df, x=0, y=1, color = df.label.astype(str))
```



## Building the Generator Model

# Generative Adversarial Networks



Write a method to create the Generator model architecture

```
In [34]: def build_generator():
    model = Sequential()
    model.add(Dense(32, activation = "relu", input_dim = 29, kernel_initializer='he_normal'))
    model.add(BatchNormalization())
    model.add(Dense(64, activation = "relu"))
    model.add(BatchNormalization())
    model.add(Dense(128, activation = "relu"))
    model.add(BatchNormalization())

    #output later
    model.add(Dense(29, activation = "relu"))
    model.compile(optimizer = "adam", loss = "binary_crossentropy")
    model.summary()

    return model
```

```
In [35]: build_generator()
```

```
/opt/anaconda3/lib/python3.12/site-packages/keras/src/layers/core/dense.py:87: UserWarning:
```

```
Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Par
dense (Dense)	(None, 32)	
batch_normalization (BatchNormalization)	(None, 32)	
dense_1 (Dense)	(None, 64)	2
batch_normalization_1 (BatchNormalization)	(None, 64)	
dense_2 (Dense)	(None, 128)	8
batch_normalization_2 (BatchNormalization)	(None, 128)	
dense_3 (Dense)	(None, 29)	3

**Total params:** 16,029 (62.61 KB)

**Trainable params:** 15,581 (60.86 KB)

**Non-trainable params:** 448 (1.75 KB)

Out[35]: <Sequential name=sequential, built=True>

## Building the Discriminator Model

Write a method to create the Discriminator model architecture

```
In [38]: def build_discriminator():
model = Sequential()
model.add(Dense(128, activation = "relu", input_dim = 29, kernel_initializer='he_normal'))
model.add(Dense(64, activation = "relu"))
model.add(Dense(32, activation = "relu"))
model.add(Dense(32, activation = "relu"))
model.add(Dense(16, activation = "relu"))
model.add(Dense(1, activation = "sigmoid"))
model.compile(optimizer = "adam", loss = "binary_crossentropy")
model.summary()
return model
```

```
In [39]: build_discriminator()
```

/opt/anaconda3/lib/python3.12/site-packages/keras/src/layers/core/dense.py:87: UserWarning:

Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

**Model: "sequential\_1"**

Layer (type)	Output Shape	Par
dense_4 (Dense)	(None, 128)	3
dense_5 (Dense)	(None, 64)	8
dense_6 (Dense)	(None, 32)	2
dense_7 (Dense)	(None, 32)	1
dense_8 (Dense)	(None, 16)	
dense_9 (Dense)	(None, 1)	

**Total params:** 15,777 (61.63 KB)

**Trainable params:** 15,777 (61.63 KB)

**Non-trainable params:** 0 (0.00 B)

Out[39]: <Sequential name=sequential\_1, built=True>

## Combine Generator and Discriminator Models to Build the GAN

```
In [41]: def build_GAN(generator, discriminator):
gan_input = Input(shape = (generator.input_shape[1],))
x = generator(gan_input)
gan_output = discriminator(x)
gan = Model(gan_input, gan_output)
gan.compile(optimizer="adam", loss="binary_crossentropy")
gan.summary()

# Freeze the discriminator from training
discriminator.trainable = False
return gan
```

Let's create a method that generates synthetic data using the Generator

```
In [43]: def generate_synthetic_data(generator, num_sample):
noise = np.random.normal(0, 1, (num_sample, generator.input_shape[1]))
fake_data = generator.predict(noise)
return fake_data
```

## Train and evaluate our GAN

- Defining some variables
- Creating our GAN
- Training the GAN
- Monitor the GAN performance using PCA

```
In [71]: def monitor_generator(generator):
# Initialize a PCA (Principal Component Analysis) object with 2 componen
pca = PCA(n_components=2)
```



```

# Drop the 'Class' column from the fraud dataset to get real data
real_fraud_data = data_fraudulent.drop("Class", axis=1)

# Transform the real fraud data using PCA
transformed_data_real = pca.fit_transform(real_fraud_data.values)

# Create a DataFrame for the transformed real data and add a 'label' column
df_real = pd.DataFrame(transformed_data_real)
df_real['label'] = "real"

# Generate synthetic fraud data using the provided generator and specify
synthetic_fraud_data = generate_synthetic_data(generator, 492)

# Transform the synthetic fraud data using PCA
transformed_data_fake = pca.fit_transform(synthetic_fraud_data)

# Create a DataFrame for the transformed fake data and add a 'label' column
df_fake = pd.DataFrame(transformed_data_fake)
df_fake["label"] = "fake"

# Concatenate the real and fake data DataFrames
df_combined = pd.concat([df_real, df_fake])

# Create a scatterplot to visualize the data points, using the first and
# and color points based on the 'label' column, with a size of 10
plt.figure()
sns.scatterplot(data=df_combined, x=0, y=1, hue='label', s=10)
plt.show()

```

```

In [75]: generator = build_generator()
discriminator = build_discriminator()
gan = build_GAN(generator, discriminator)
gan.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy"])

num_epochs = 100
batch_size = 64
half_batch = int(batch_size / 2)

for epoch in range(num_epochs):
    # Generate synthetic data
    X_fake = generate_synthetic_data(generator, half_batch)
    y_fake = np.zeros((half_batch, 1))

    # Sample a batch of real data
    X_real = data_fraudulent.drop("Class", axis=1).sample(n=half_batch, random_state=epoch)
    y_real = np.ones((half_batch, 1))

    discriminator.compile(optimizer="adam", loss="binary_crossentropy")

    # Train the discriminator
    discriminator.trainable = True
    discriminator.train_on_batch(X_real, y_real)
    discriminator.train_on_batch(X_fake, y_fake)

    # Train the GAN (generator part)

```

```

noise = np.random.normal(0, 1, (batch_size, 29))
gan.train_on_batch(noise, np.ones((batch_size, 1)))

if epoch%10 == 0:
    monitor_generator(generator)

```

/opt/anaconda3/lib/python3.12/site-packages/keras/src/layers/core/dense.py:87: UserWarning:

Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

Model: "sequential\_12"

Layer (type)	Output Shape	Par
dense_60 (Dense)	(None, 32)	
batch_normalization_18 (BatchNormalization)	(None, 32)	
dense_61 (Dense)	(None, 64)	2
batch_normalization_19 (BatchNormalization)	(None, 64)	
dense_62 (Dense)	(None, 128)	8
batch_normalization_20 (BatchNormalization)	(None, 128)	
dense_63 (Dense)	(None, 29)	3

Total params: 16,029 (62.61 KB)

Trainable params: 15,581 (60.86 KB)

Non-trainable params: 448 (1.75 KB)

Model: "sequential\_13"

Layer (type)	Output Shape	Par
dense_64 (Dense)	(None, 128)	3
dense_65 (Dense)	(None, 64)	8
dense_66 (Dense)	(None, 32)	2
dense_67 (Dense)	(None, 32)	1
dense_68 (Dense)	(None, 16)	
dense_69 (Dense)	(None, 1)	

Total params: 15,777 (61.63 KB)

Trainable params: 15,777 (61.63 KB)

Non-trainable params: 0 (0.00 B)

Model: "functional\_96"

Layer (type)	Output Shape	Par
input_layer_19 (InputLayer)	(None, 29)	
sequential_12 (Sequential)	(None, 29)	16
sequential_13 (Sequential)	(None, 1)	15

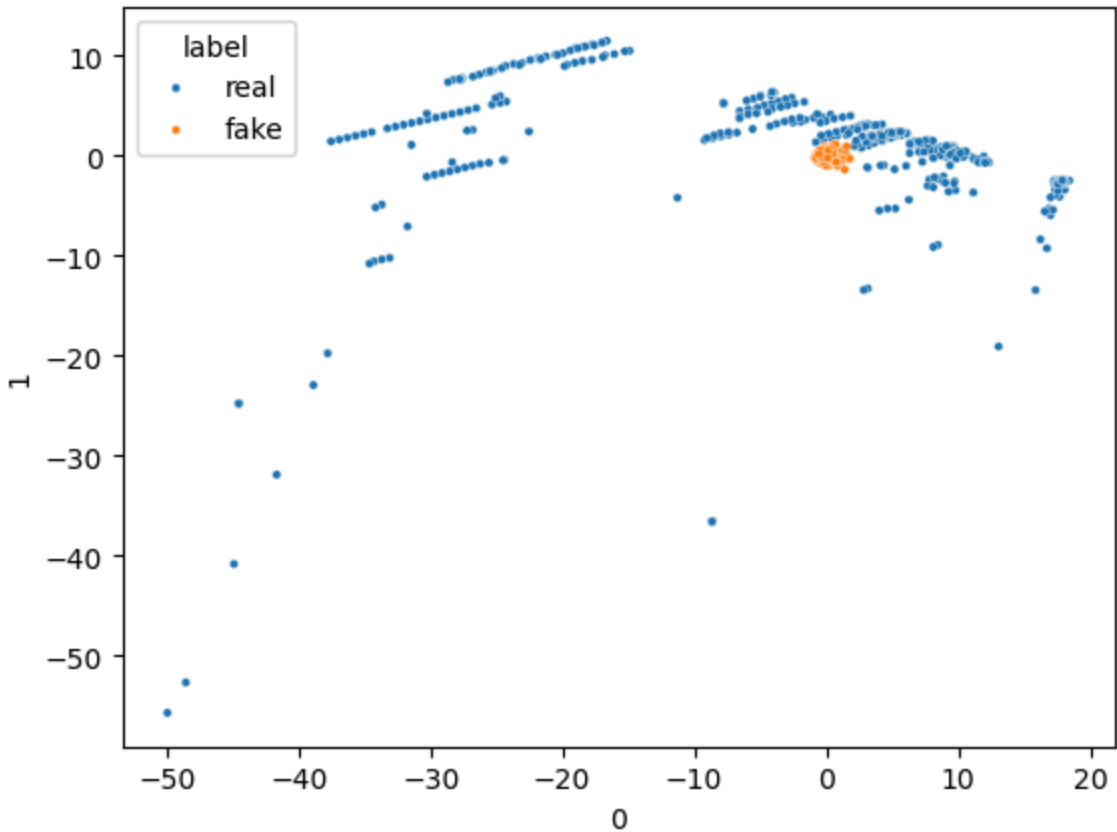
Total params: 31,806 (124.24 KB)

Trainable params: 31,358 (122.49 KB)

Non-trainable params: 448 (1.75 KB)

1/1 0s 32ms/step

16/16 0s 3ms/step



1/1 0s 7ms/step

1/1 0s 5ms/step

1/1 0s 4ms/step

1/1 0s 4ms/step

1/1 0s 5ms/step

1/1 0s 5ms/step

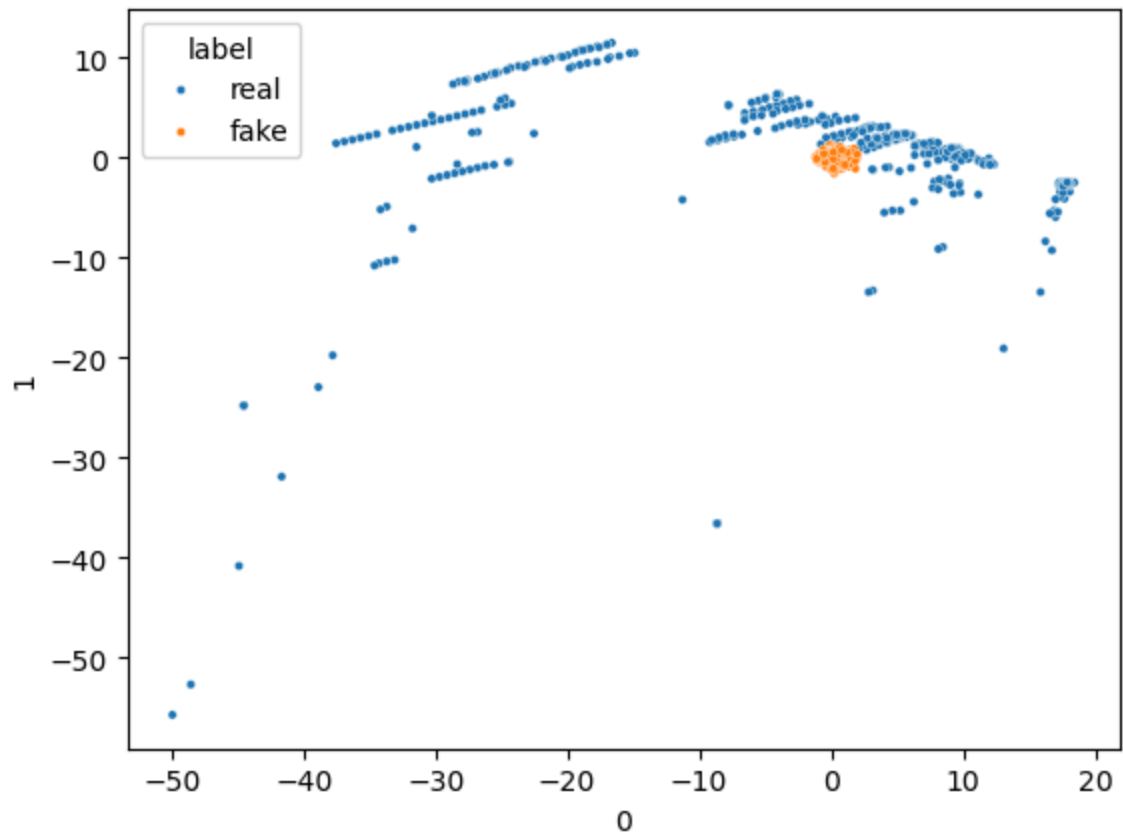
1/1 0s 5ms/step

1/1 0s 4ms/step

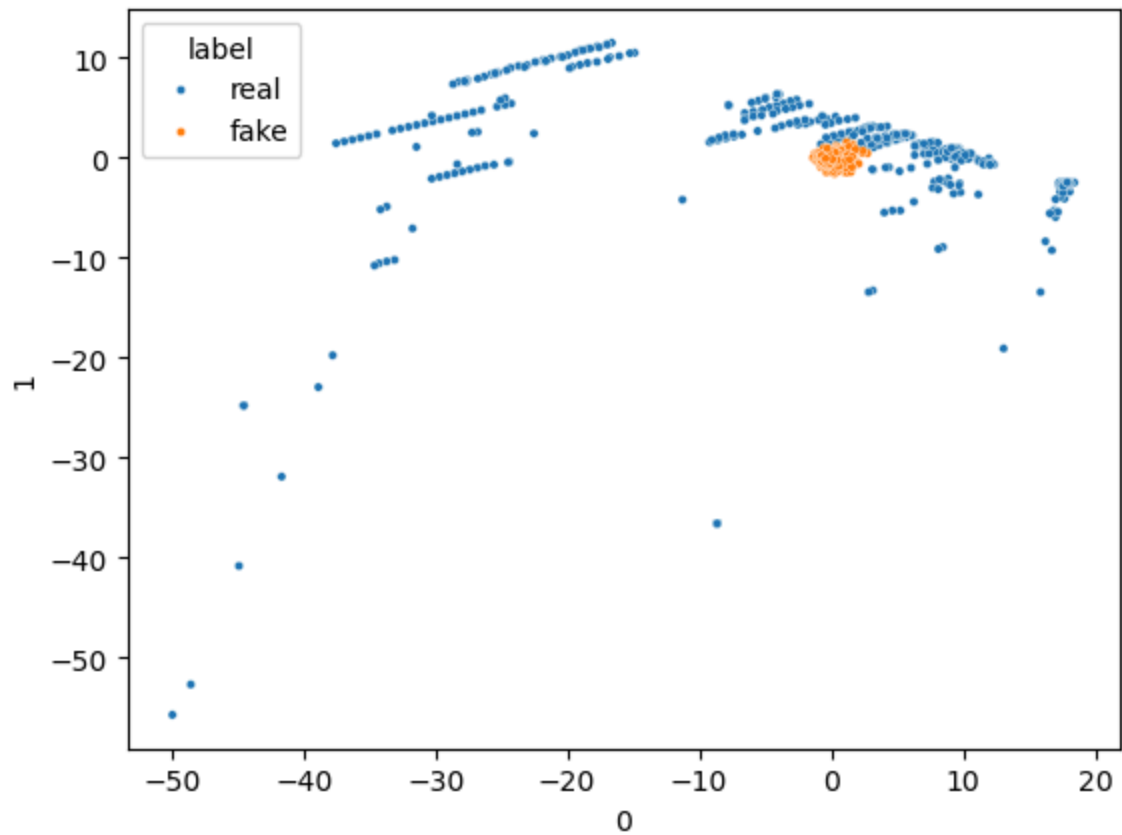
1/1 0s 5ms/step

1/1 0s 5ms/step

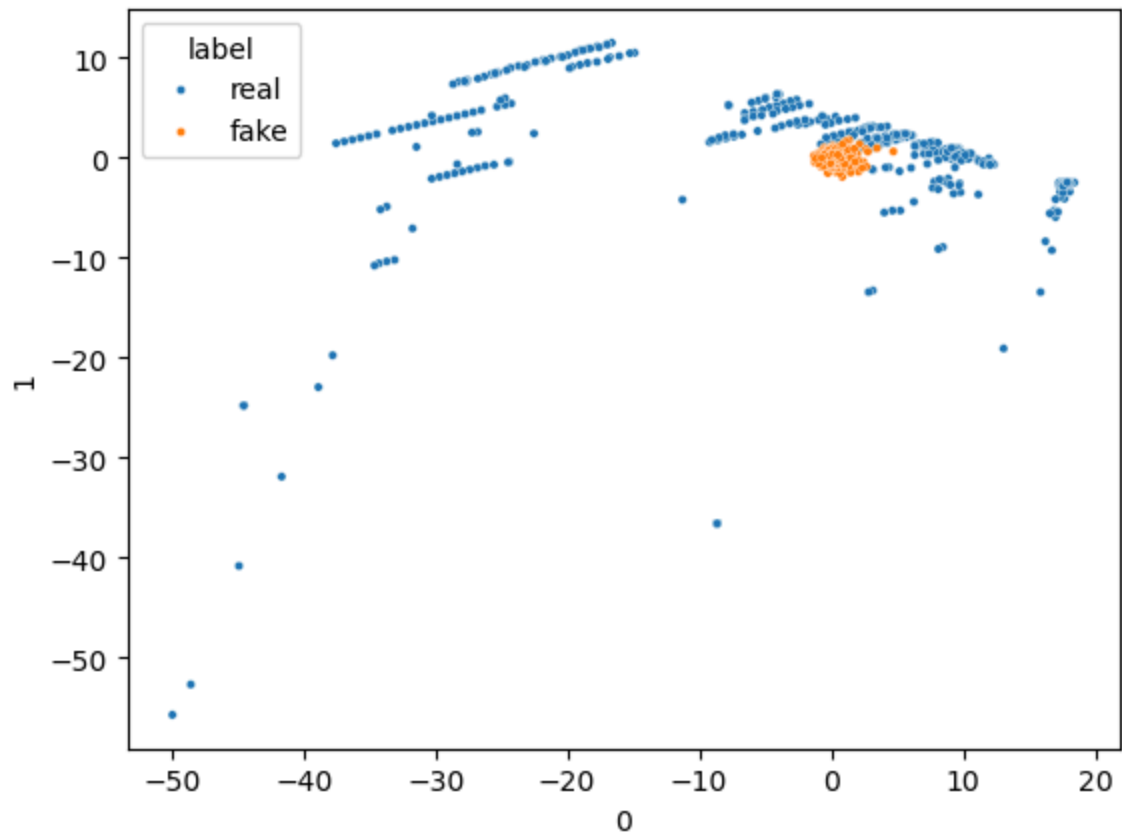
16/16 0s 666us/step



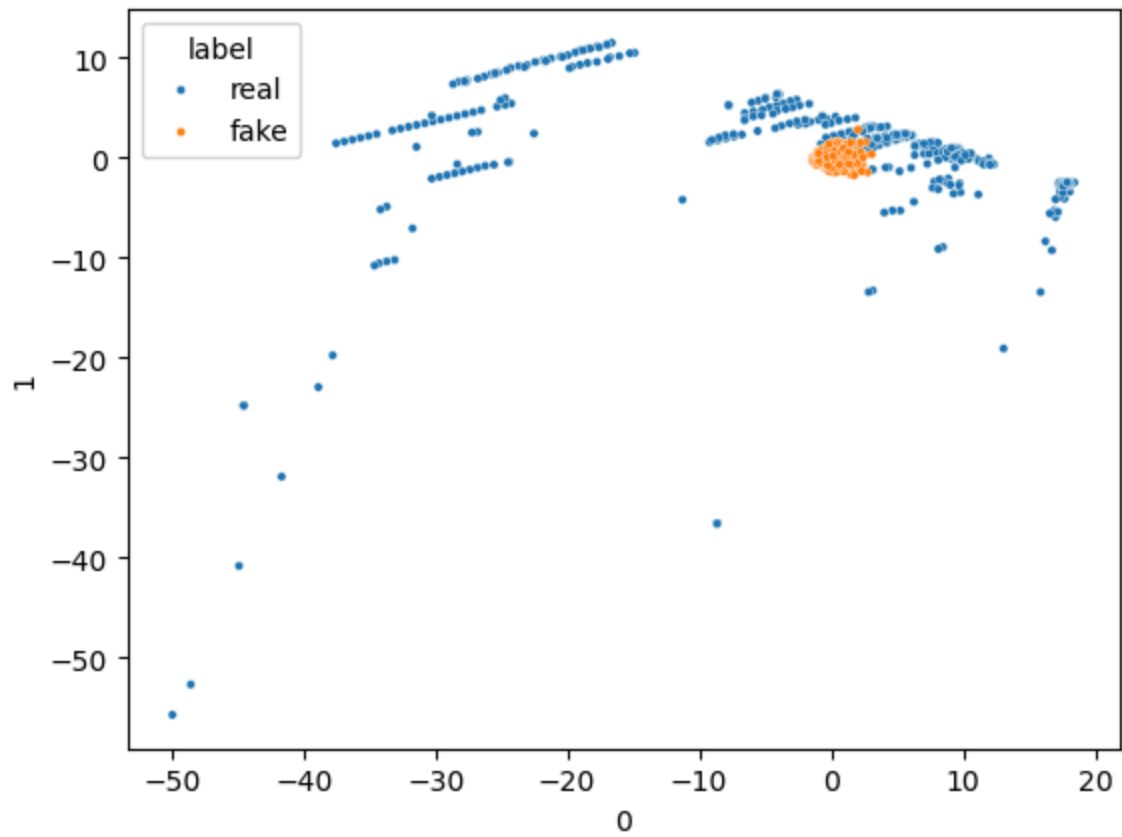
1/1	—————	0s 7ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 6ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 6ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 6ms/step
1/1	—————	0s 5ms/step
16/16	—————	0s 869us/step



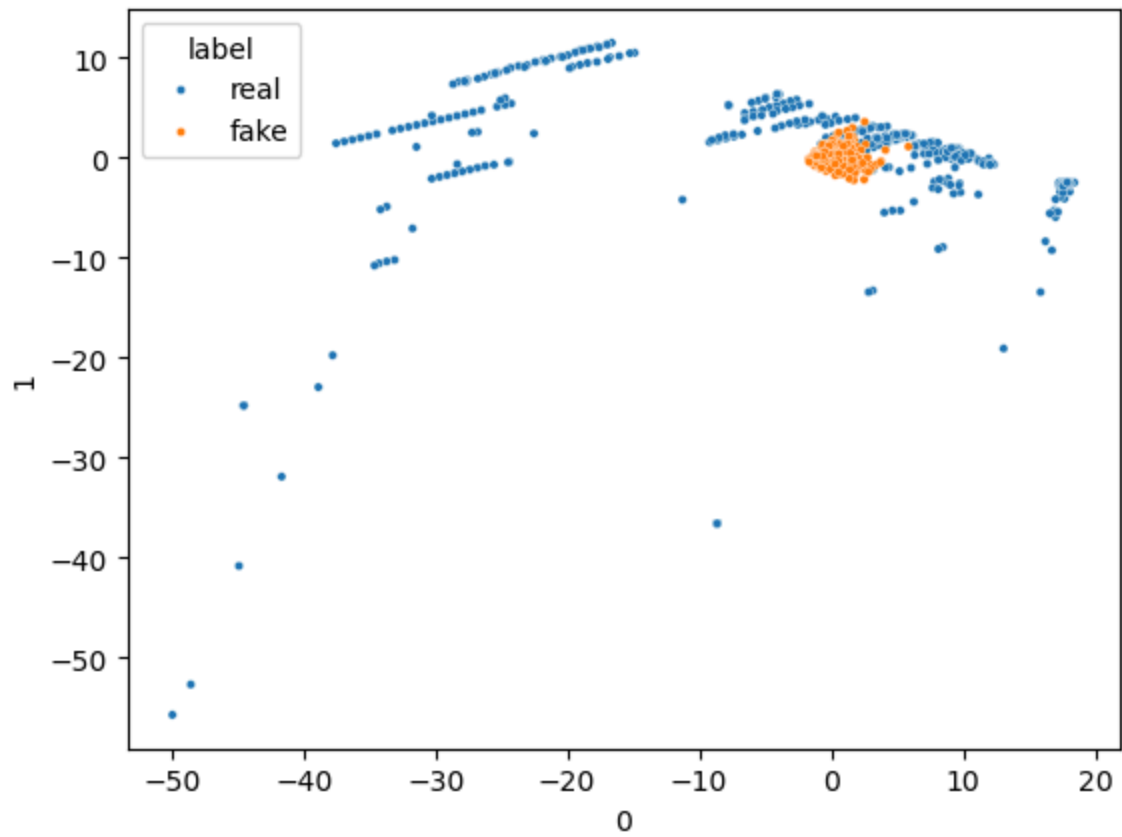
1/1	—————	0s	6ms/step
1/1	—————	0s	5ms/step
1/1	—————	0s	5ms/step
1/1	—————	0s	4ms/step
1/1	—————	0s	5ms/step
1/1	—————	0s	5ms/step
1/1	—————	0s	5ms/step
1/1	—————	0s	5ms/step
1/1	—————	0s	5ms/step
1/1	—————	0s	6ms/step
16/16	—————	0s	596us/step



1/1	—————	0s 7ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 7ms/step
16/16	—————	0s 585us/step

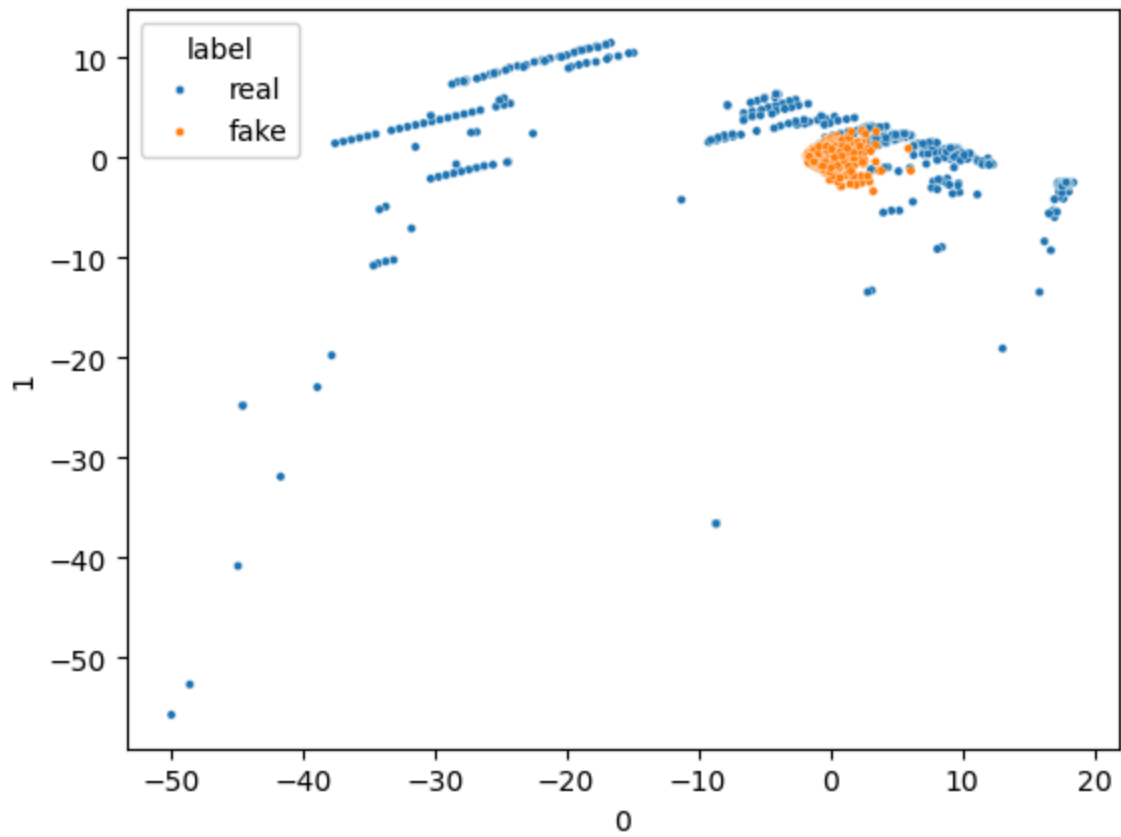


1/1	—————	0s 7ms/step
1/1	—————	0s 6ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 22ms/step
1/1	—————	0s 5ms/step
16/16	—————	0s 616us/step

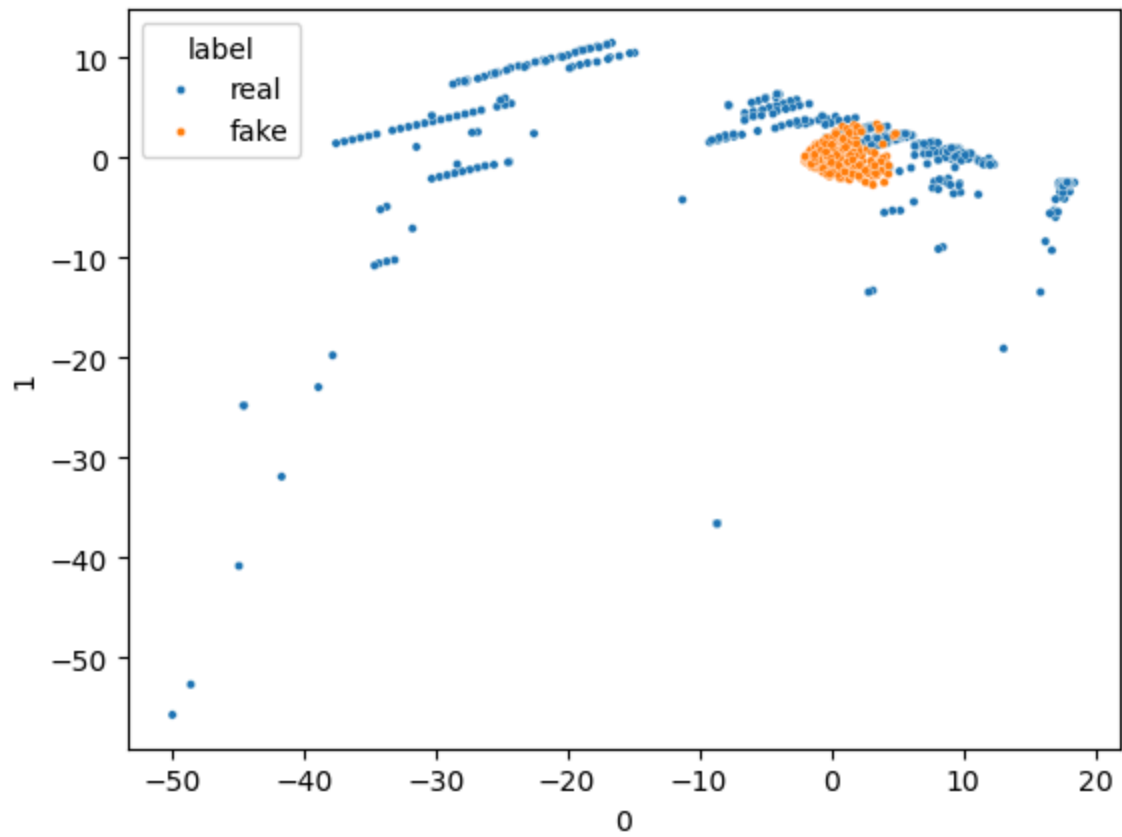


1/1	—————	0s 6ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 6ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
16/16	—————	0s 1ms/step

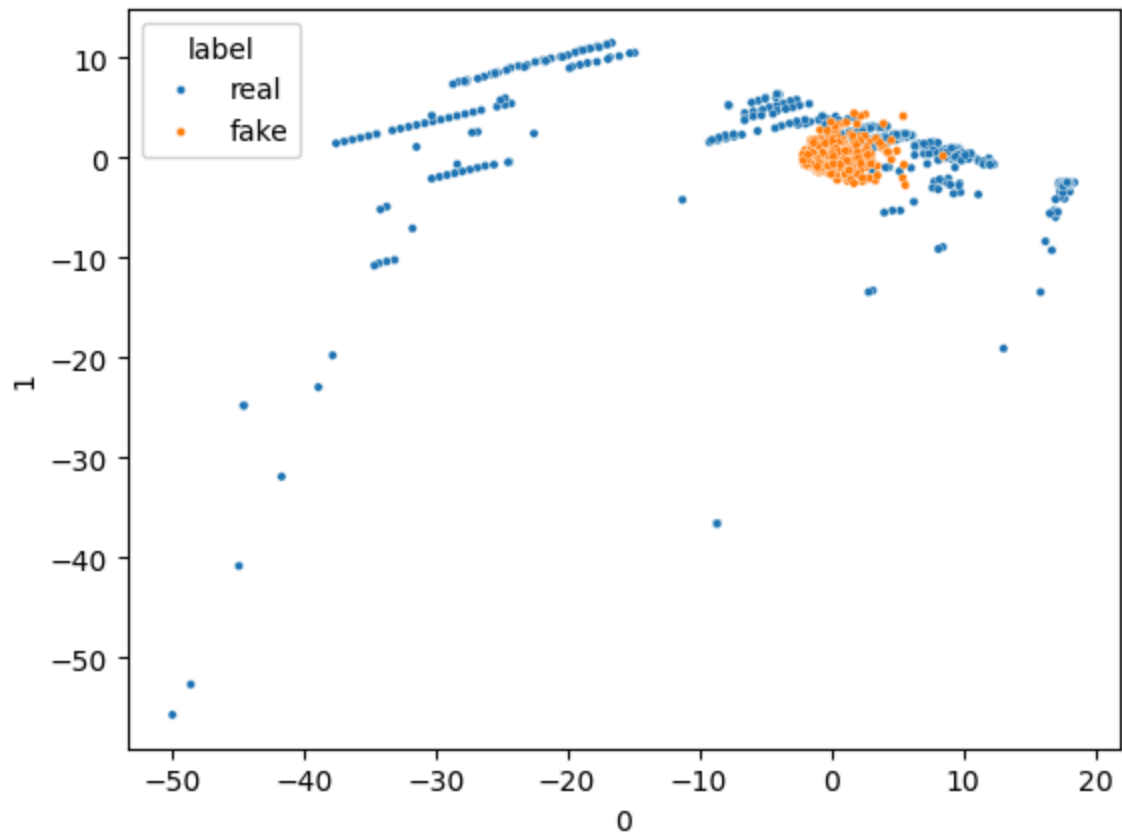




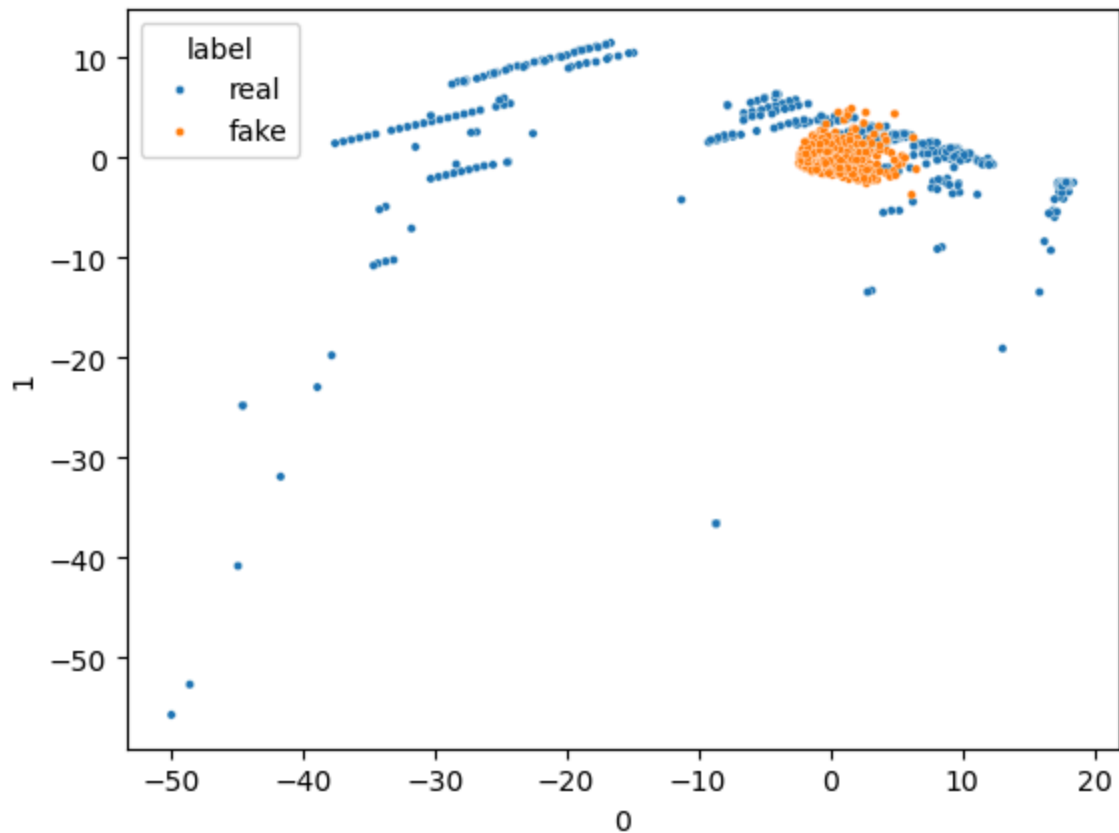
1/1	—————	0s 11ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 6ms/step
1/1	—————	0s 6ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
16/16	—————	0s 606us/step



1/1	—————	0s 8ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
16/16	—————	0s 645us/step



1/1	—————	0s 10ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 5ms/step
1/1	—————	0s 4ms/step
1/1	—————	0s 5ms/step
16/16	—————	0s 619us/step



```
1/1 ————— 0s 8ms/step
1/1 ————— 0s 4ms/step
1/1 ————— 0s 5ms/step
1/1 ————— 0s 5ms/step
1/1 ————— 0s 4ms/step
1/1 ————— 0s 4ms/step
1/1 ————— 0s 4ms/step
1/1 ————— 0s 5ms/step
1/1 ————— 0s 4ms/step
1/1 ————— 0s 4ms/step
```

## Generate synthetic data using the trained Generator

- Generate 1000 fraudulent data points using the trained generator
- Compare the distribution of real and synthetic fraudulent data points.

```
In [63]: synthetic_data = generate_synthetic_data(generator, 1000)
df = pd.DataFrame(synthetic_data)
df["label"] = "fake"

df2 = data_fraudulent.drop("Class", axis = 1)
df2["label"] = "real"
df2.columns = df.columns

combined_df = pd.concat([df, df2])
```

```
32/32 ————— 0s 263us/step
```

```
In [65]: combined_df
```

Out [65]:

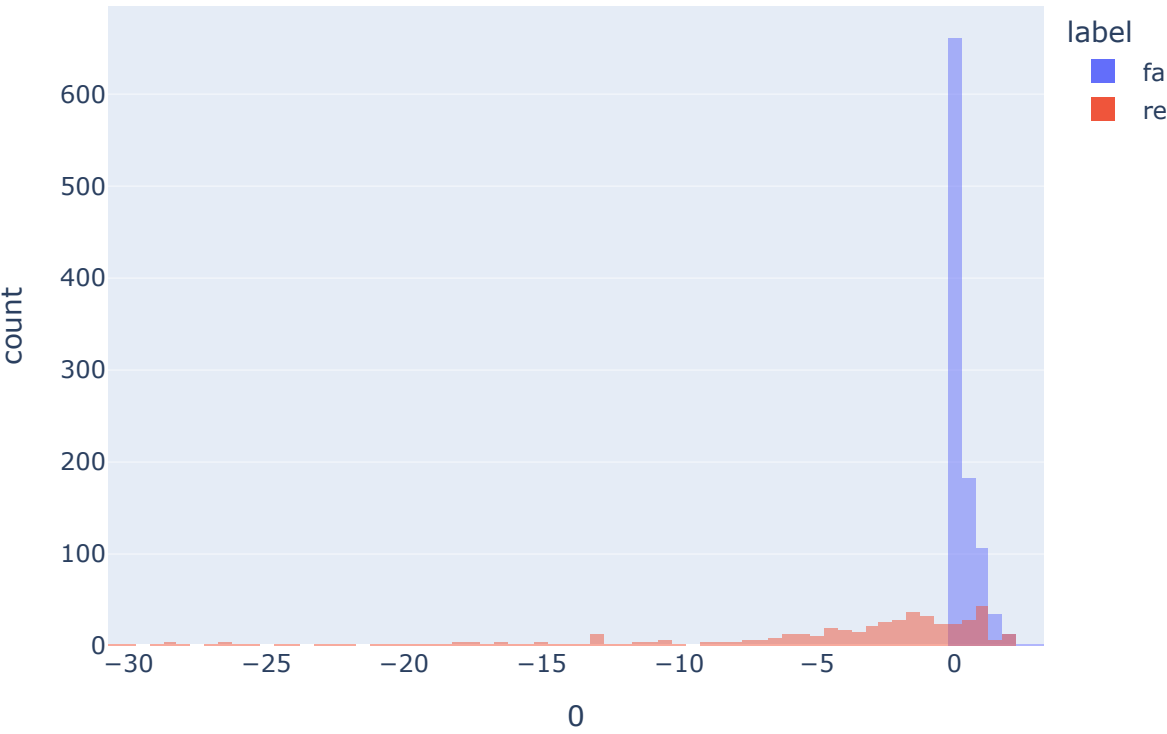
	0	1	2	3	4	5	6
0	0.488779	0.000000	0.000000	0.000000	0.192821	0.000000	0.683964
1	0.088744	0.291647	0.000000	0.000000	0.000000	0.000000	0.003341
2	0.000000	0.000000	1.074494	0.325012	0.920057	1.140243	0.000000
3	0.000000	0.000000	0.414872	0.000000	0.224516	0.000000	0.000000
4	1.341631	0.000000	0.000000	0.000000	0.058641	2.130062	0.000000
...	...	...	...	...	...	...	...
50487	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850
50488	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170
50489	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739
50490	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002
50491	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050

1492 rows x 30 columns

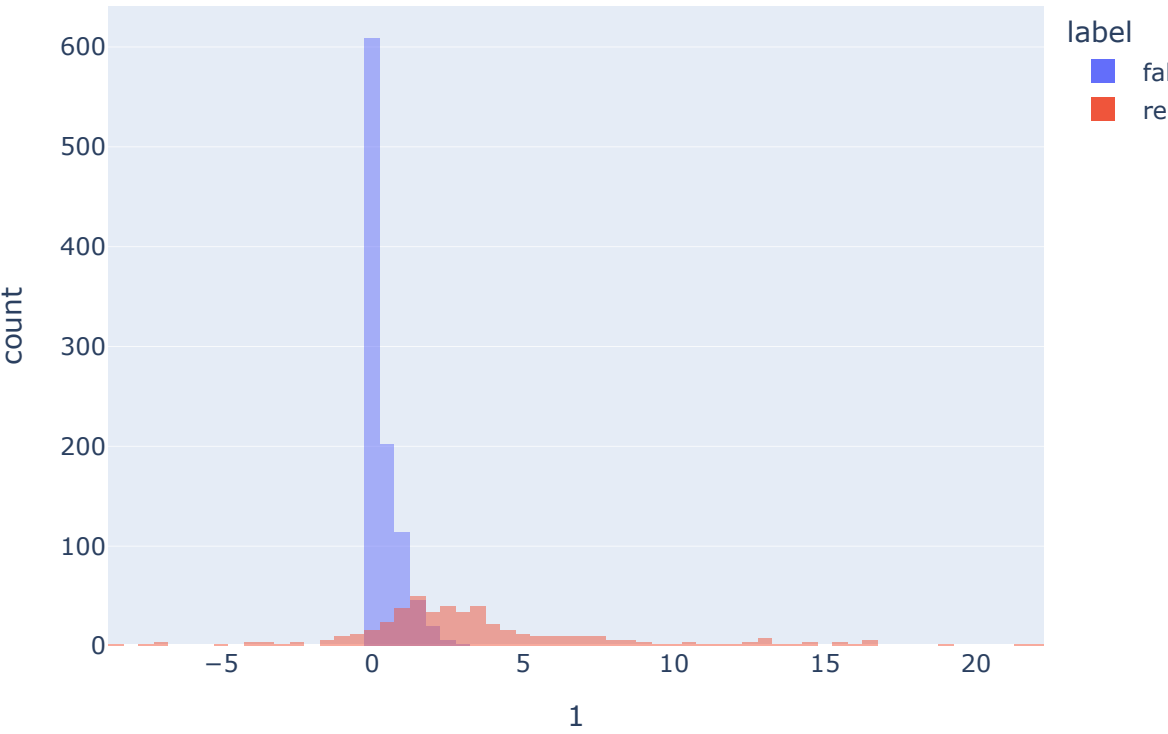
Checking the individual feature distribution of synthetic and real fraudulent data.

```
In [67]: for col in combined_df.columns:
          plt.figure()
          fig = px.histogram(combined_df, color = 'label', x=col, barmode="overlay",
                             fig.show()
```

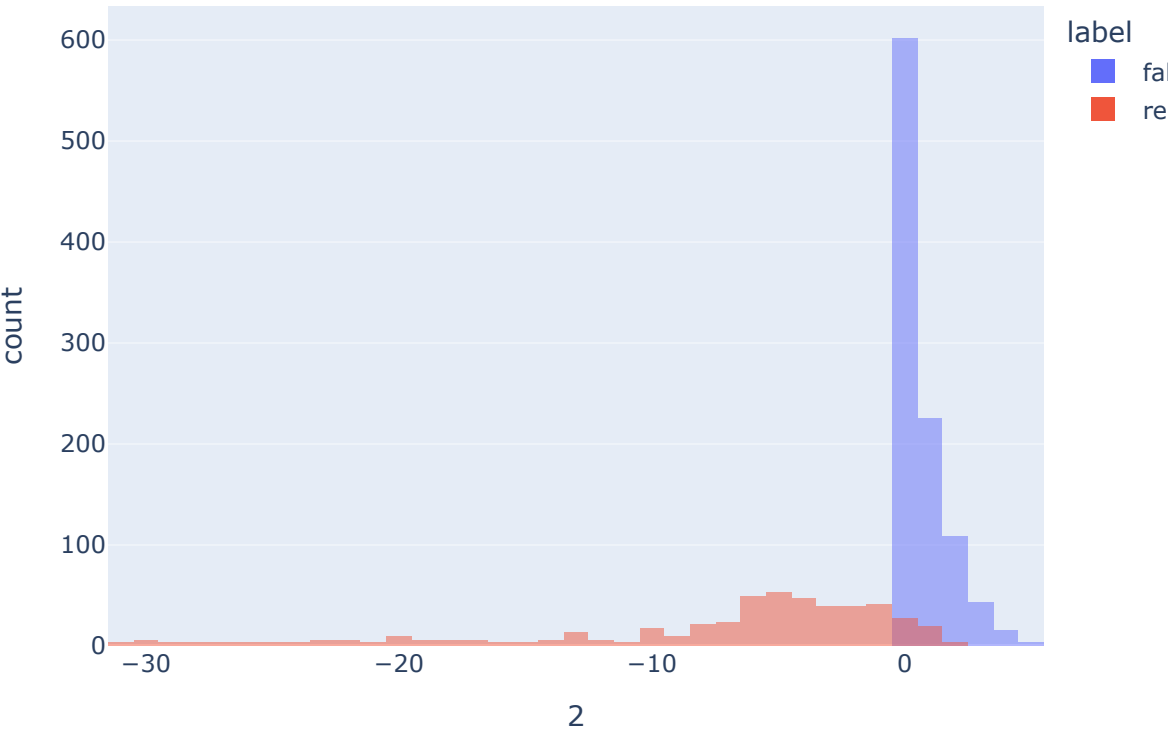
Feature 0



Feature 1

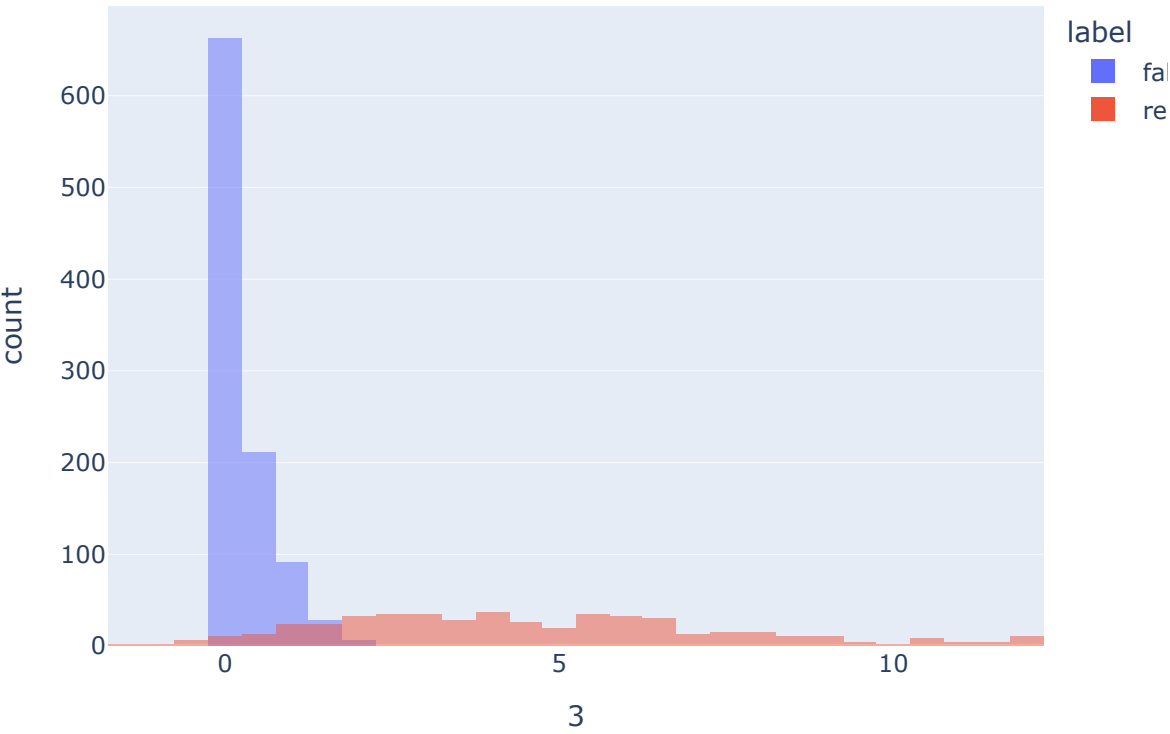


Feature 2

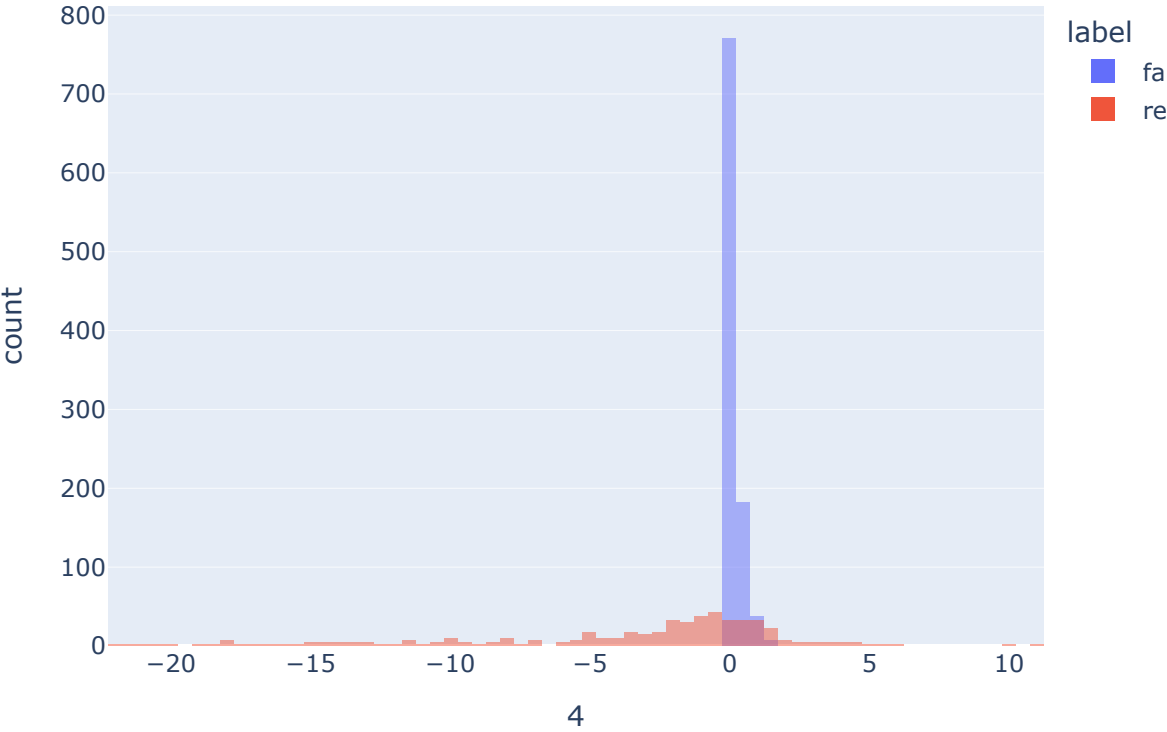




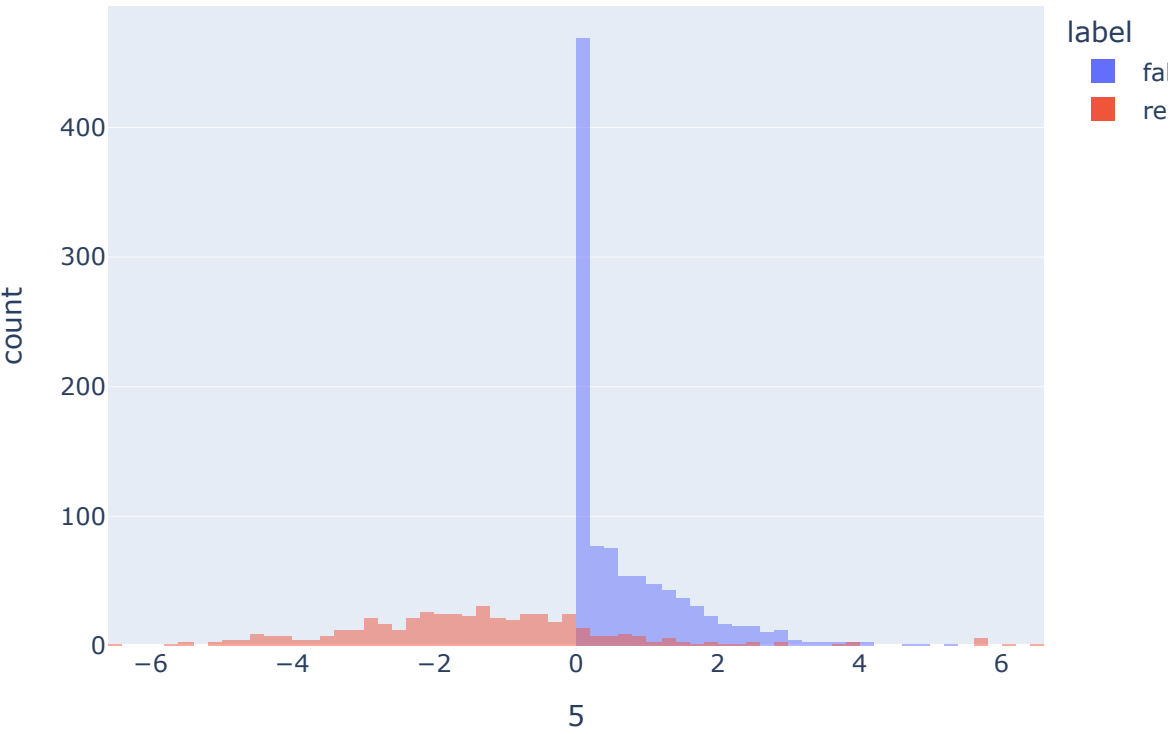
Feature 3



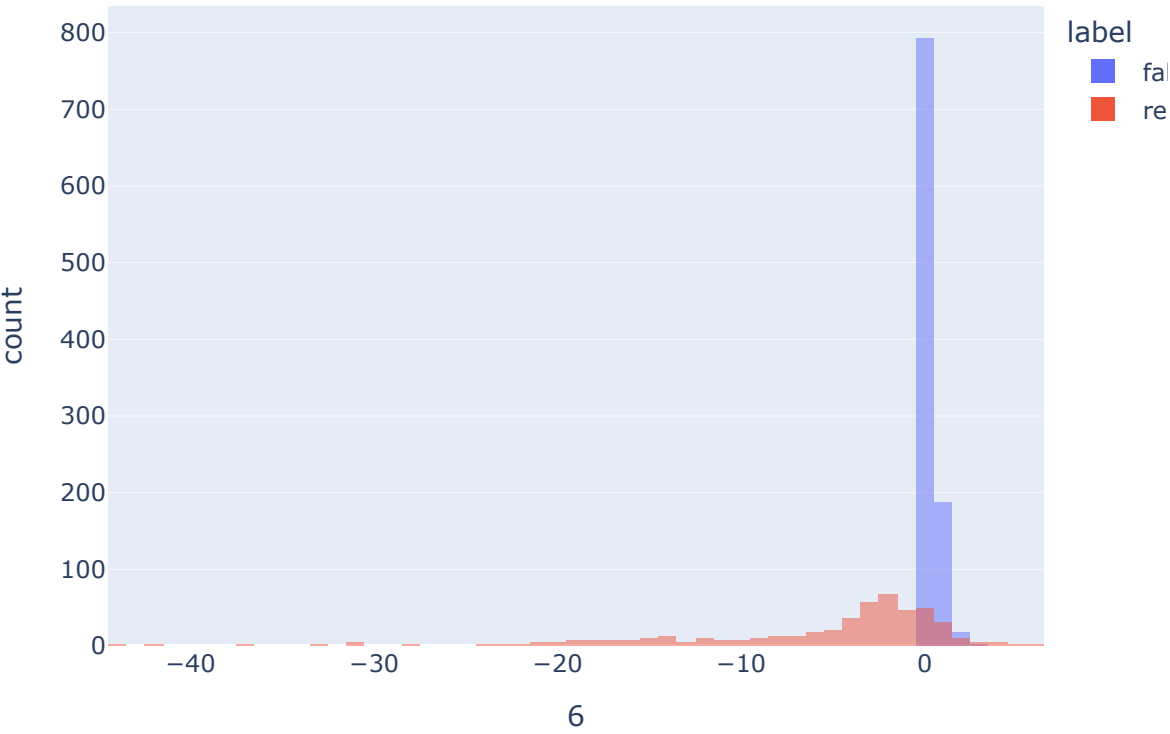
Feature 4



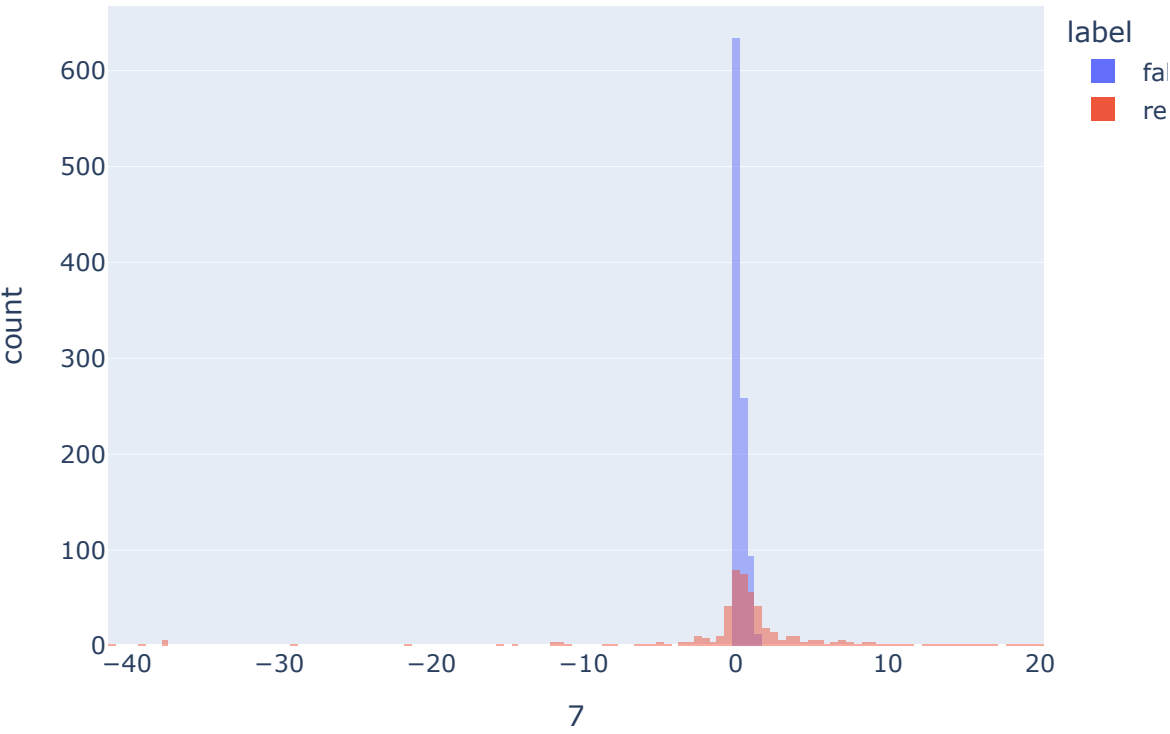
Feature 5



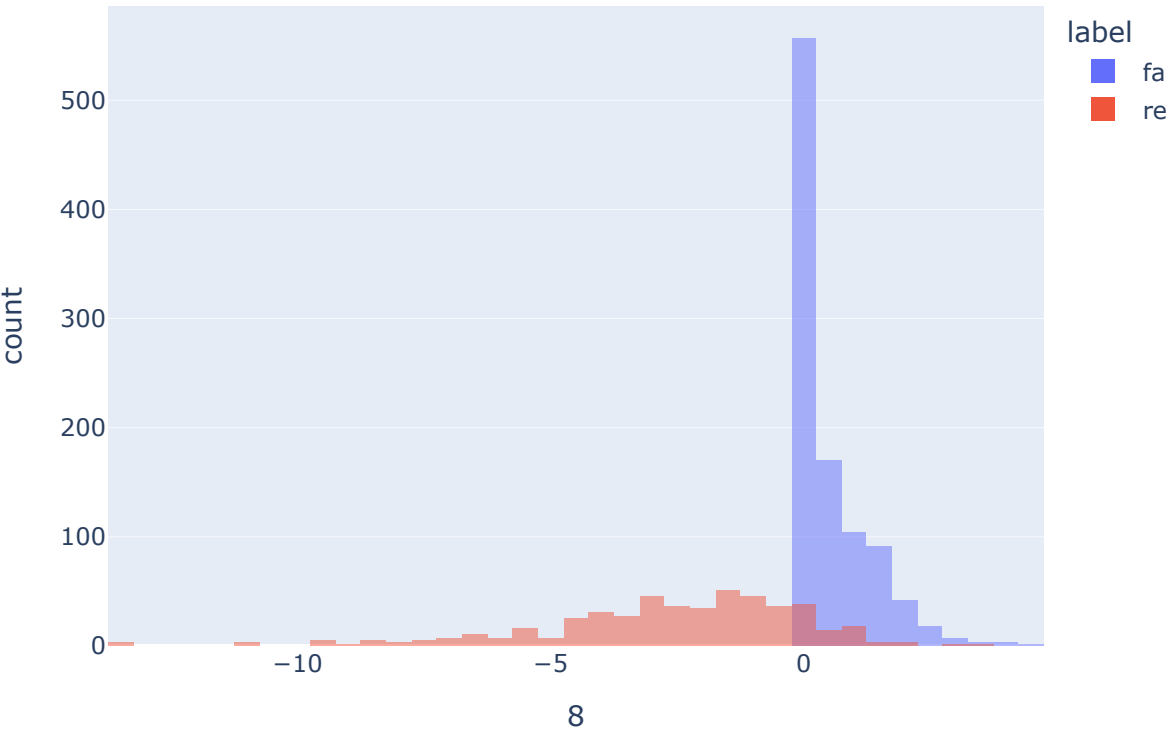
Feature 6



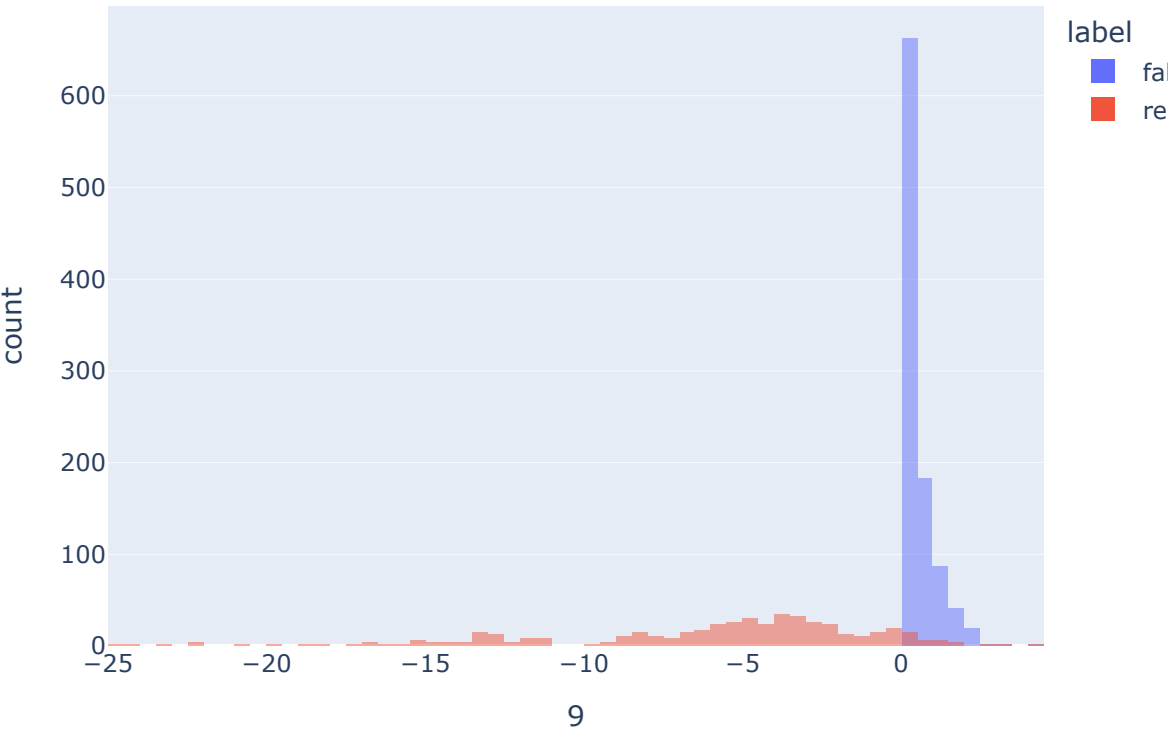
Feature 7



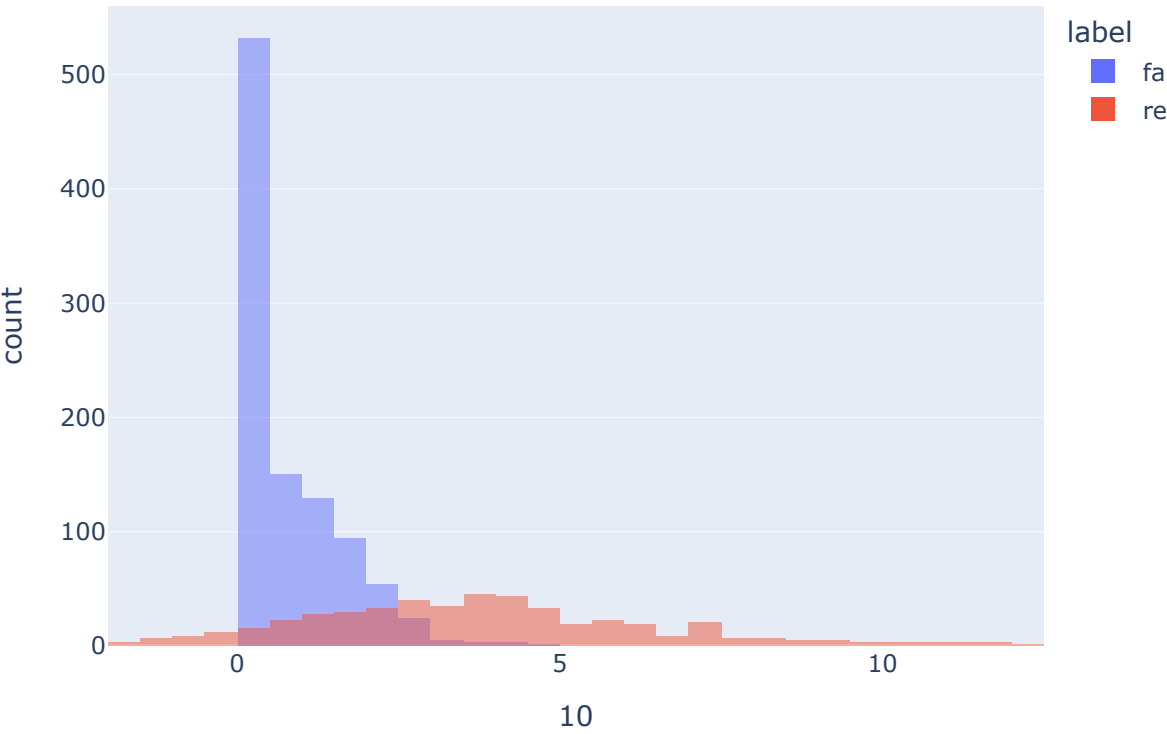
Feature 8



Feature 9

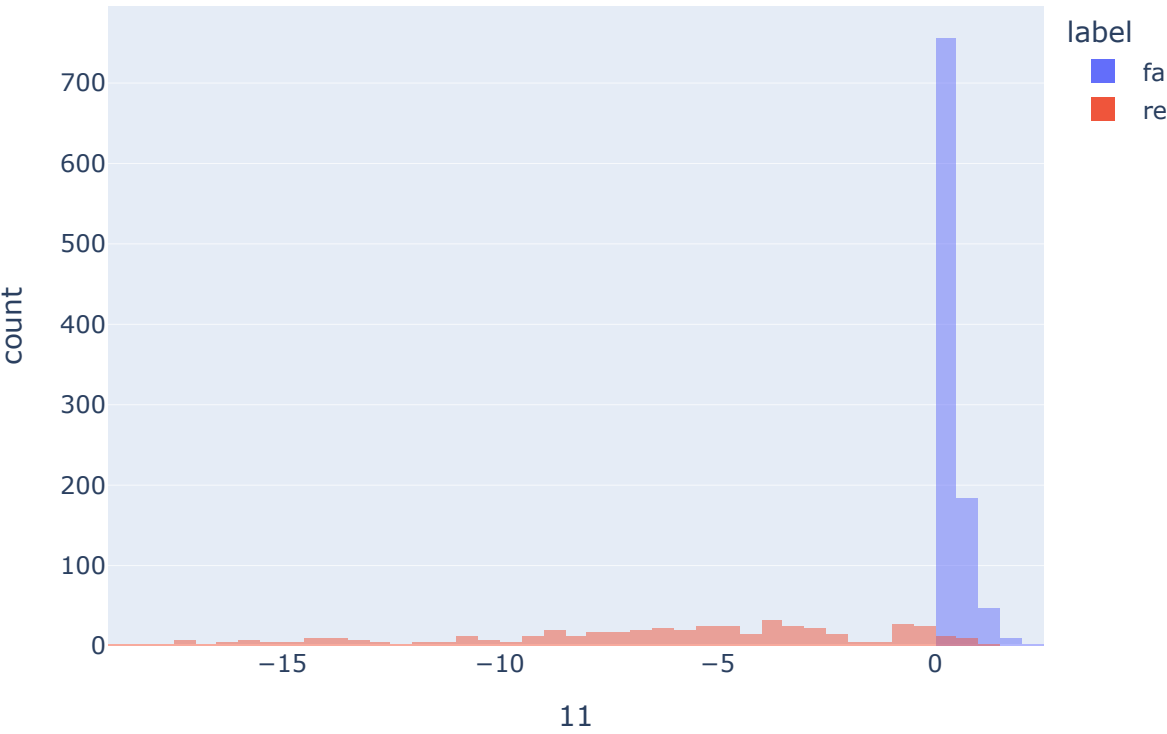


Feature 10

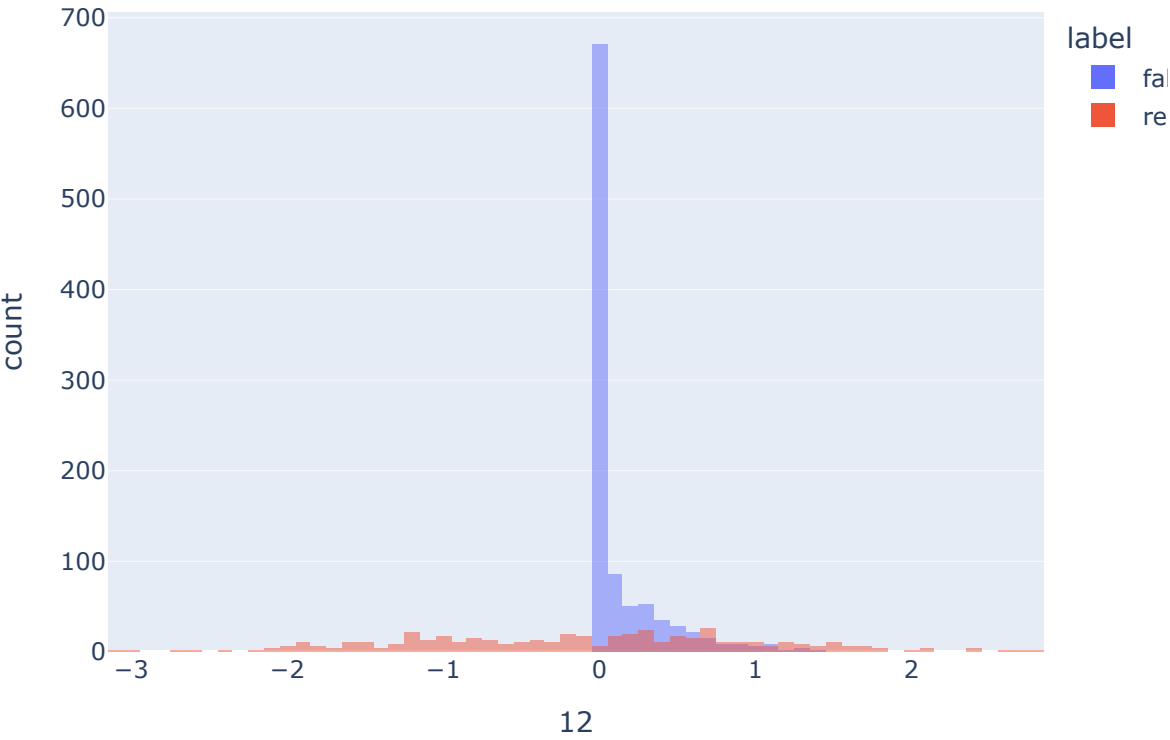




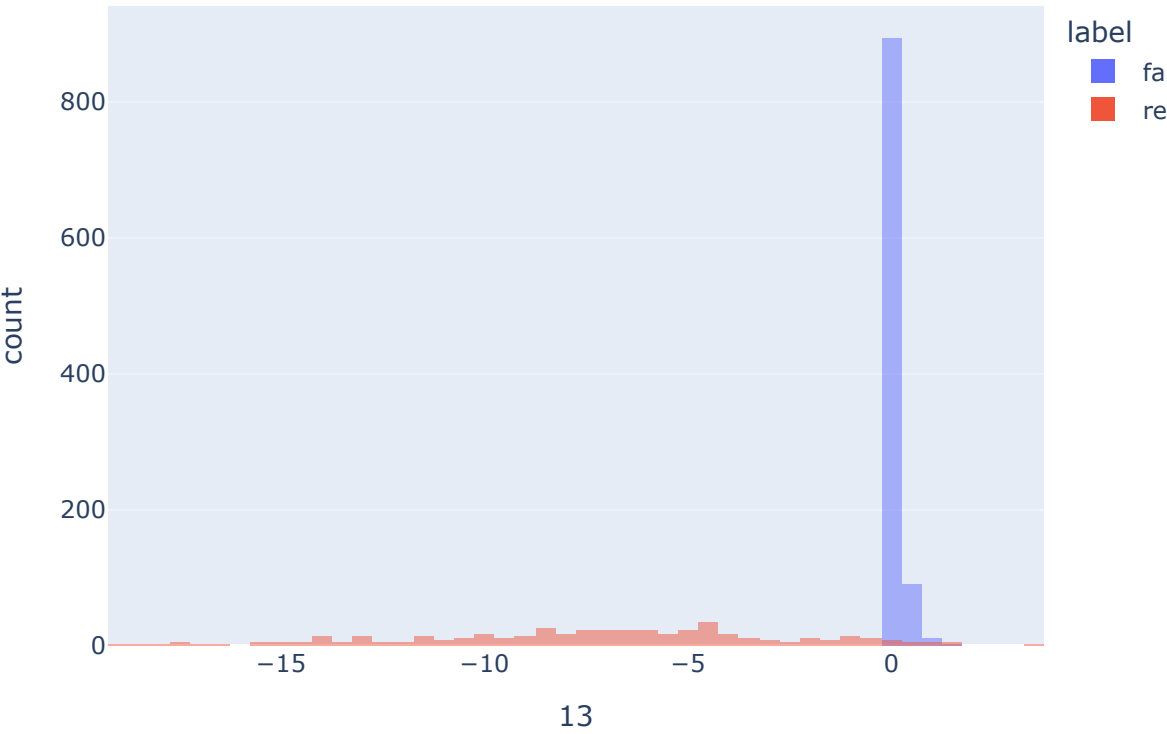
Feature 11



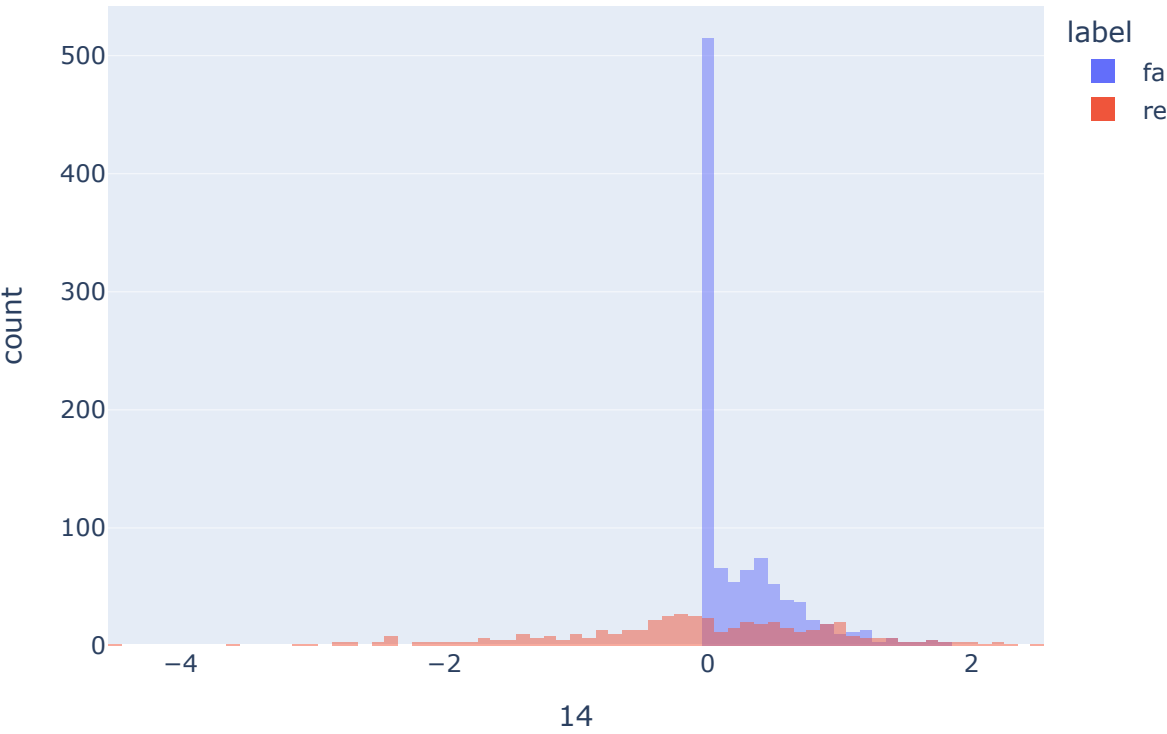
Feature 12



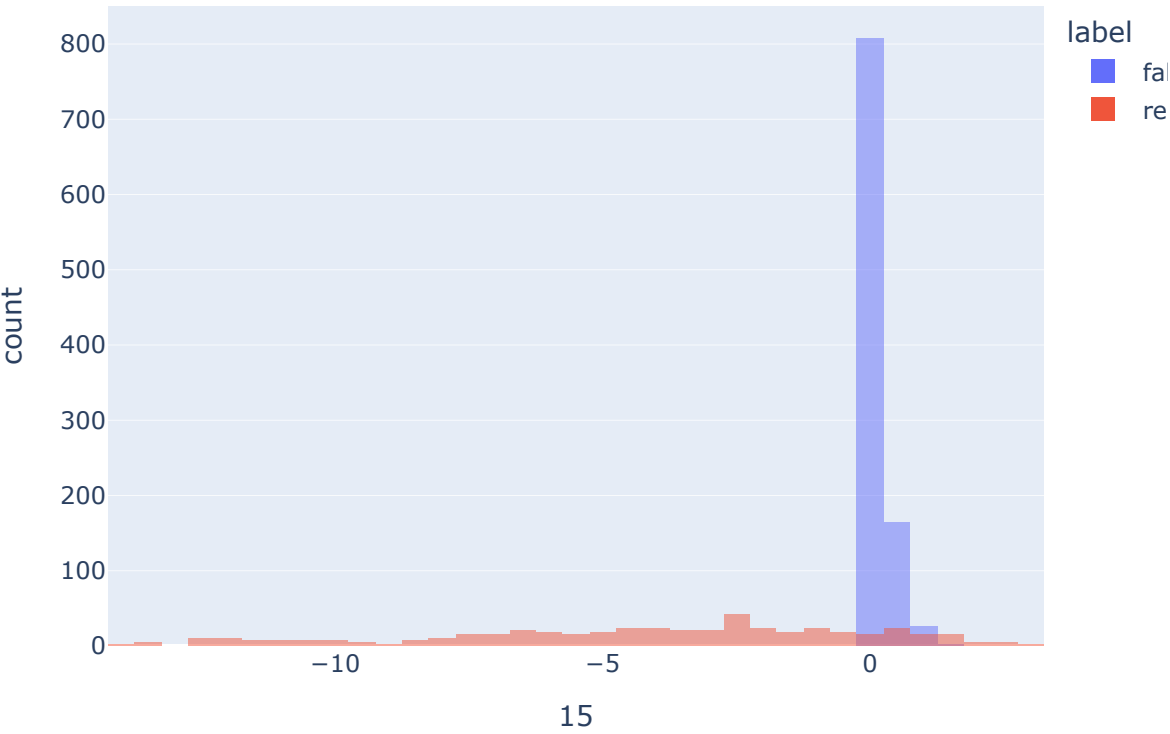
Feature 13



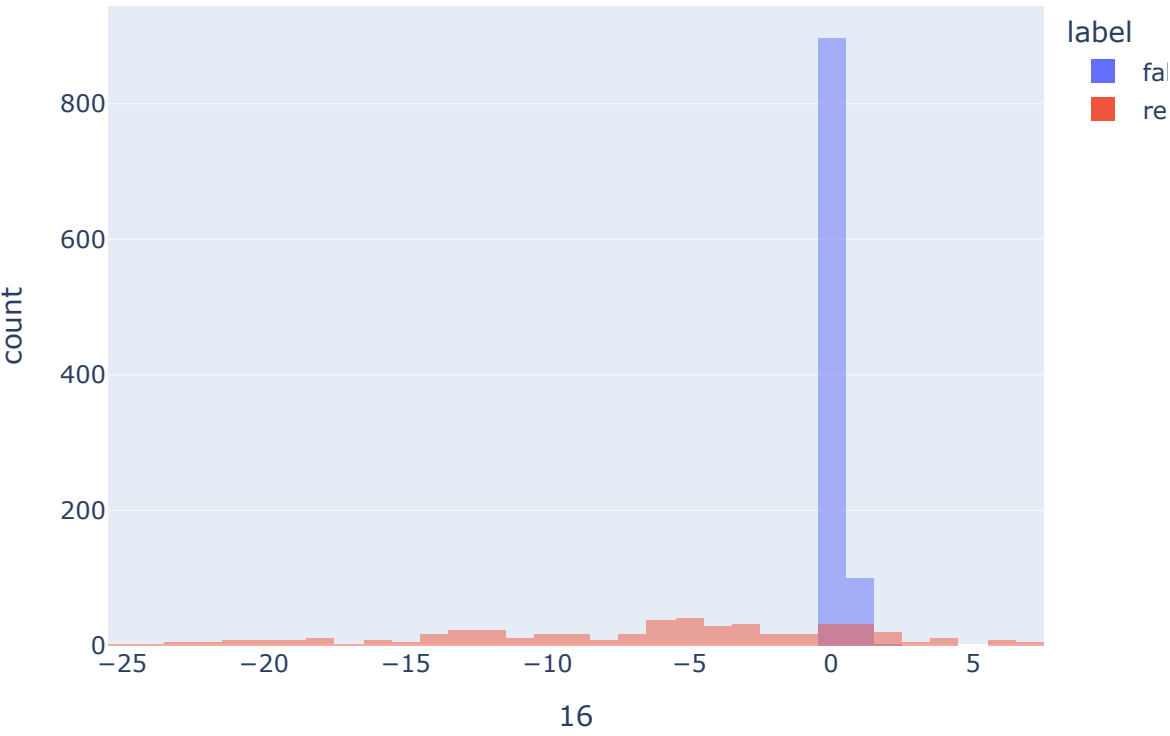
Feature 14



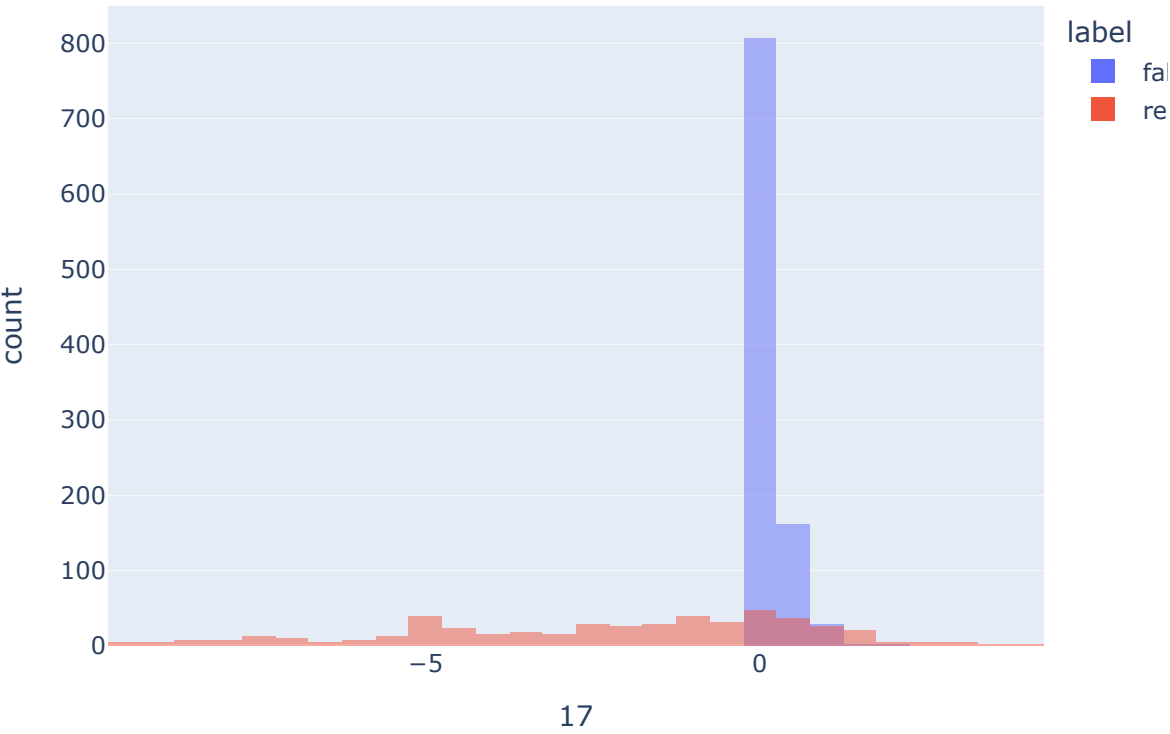
Feature 15



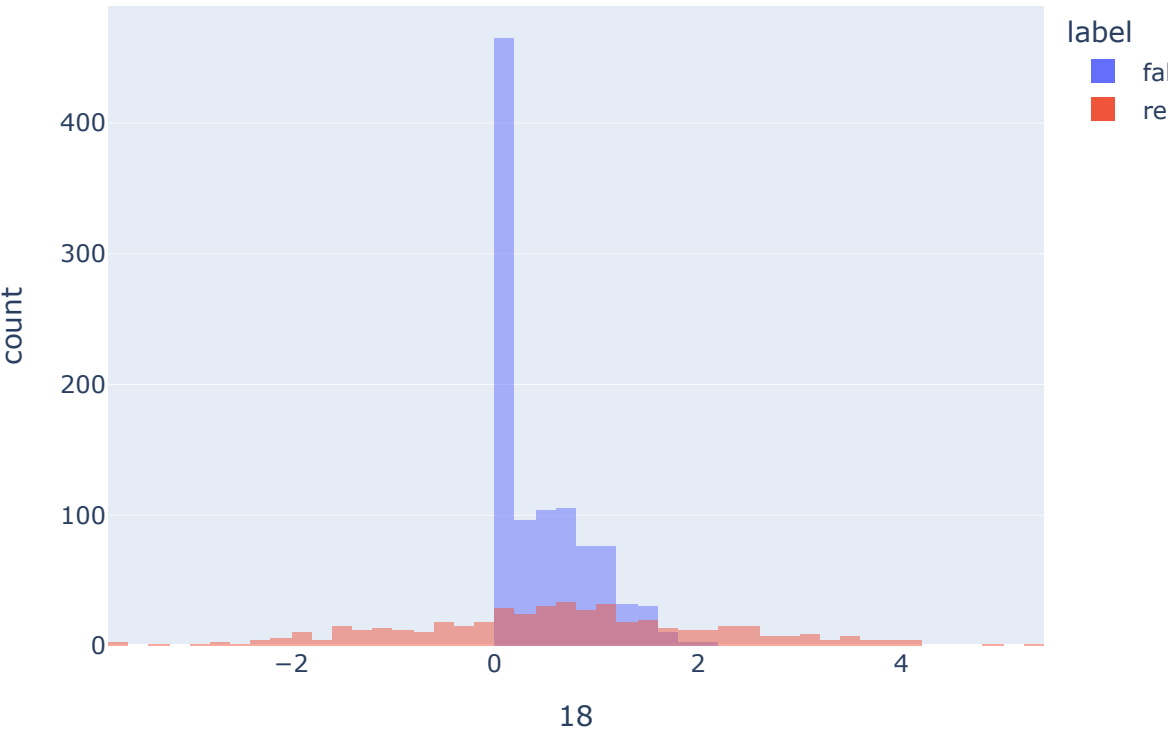
Feature 16



Feature 17

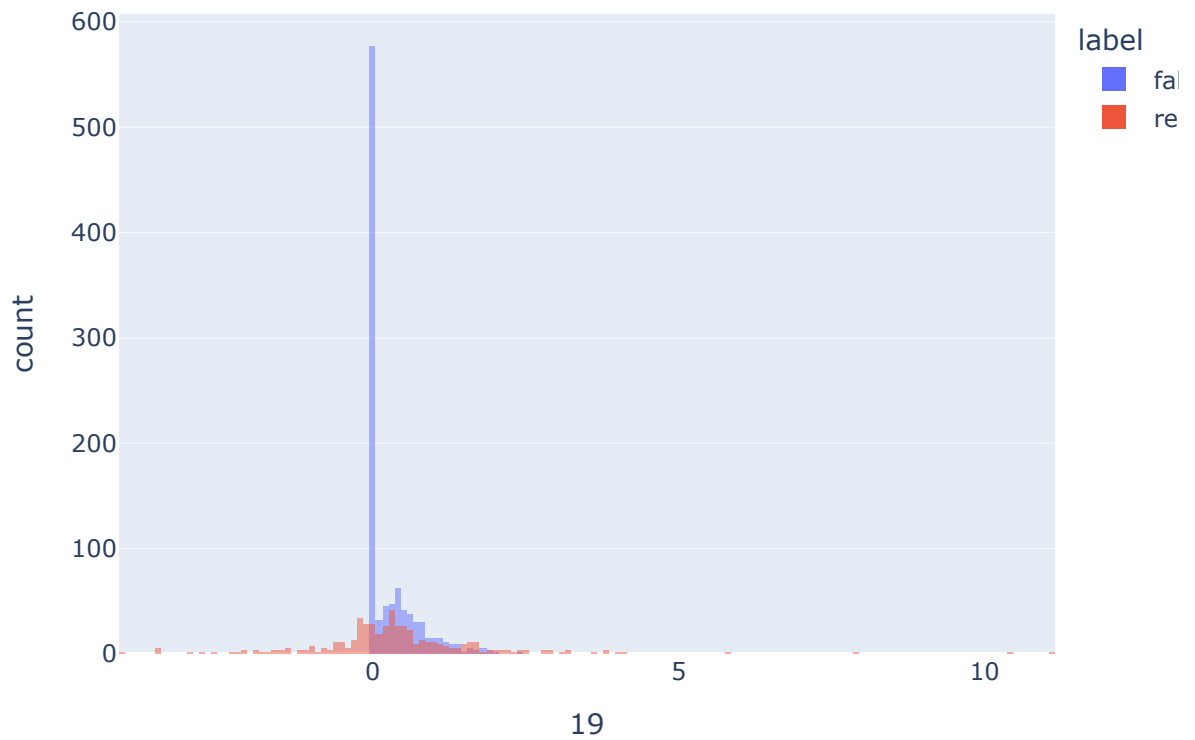


Feature 18





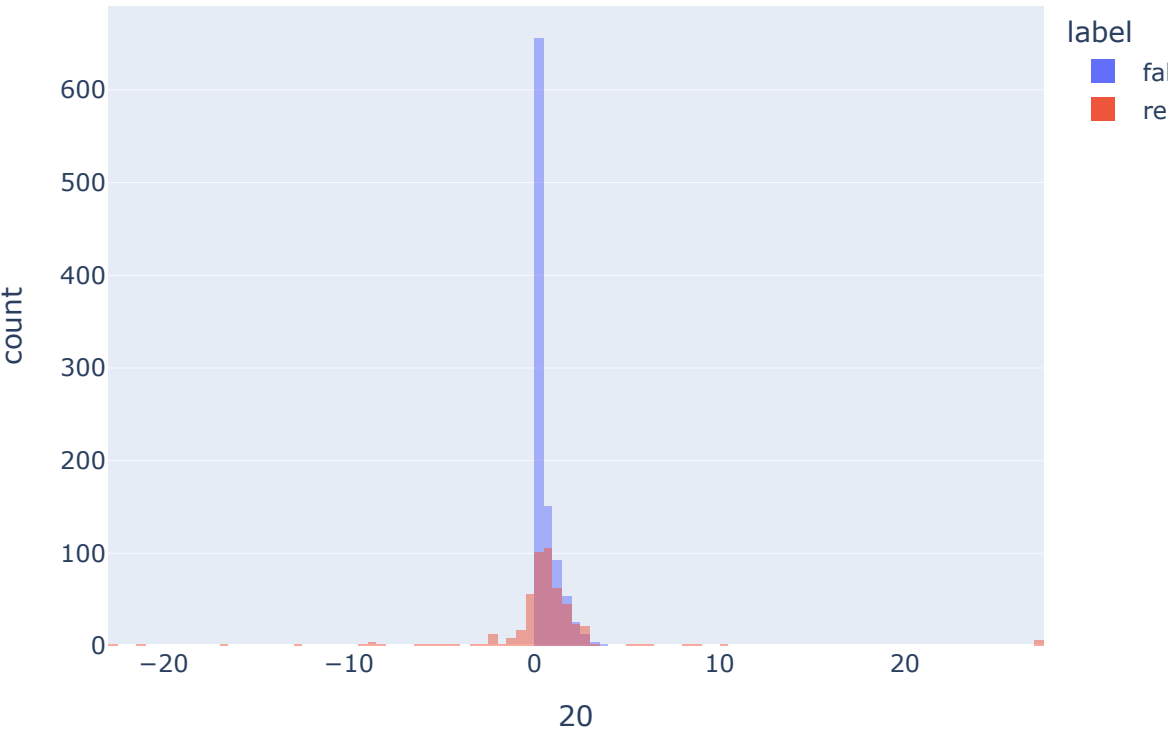
## Feature 19



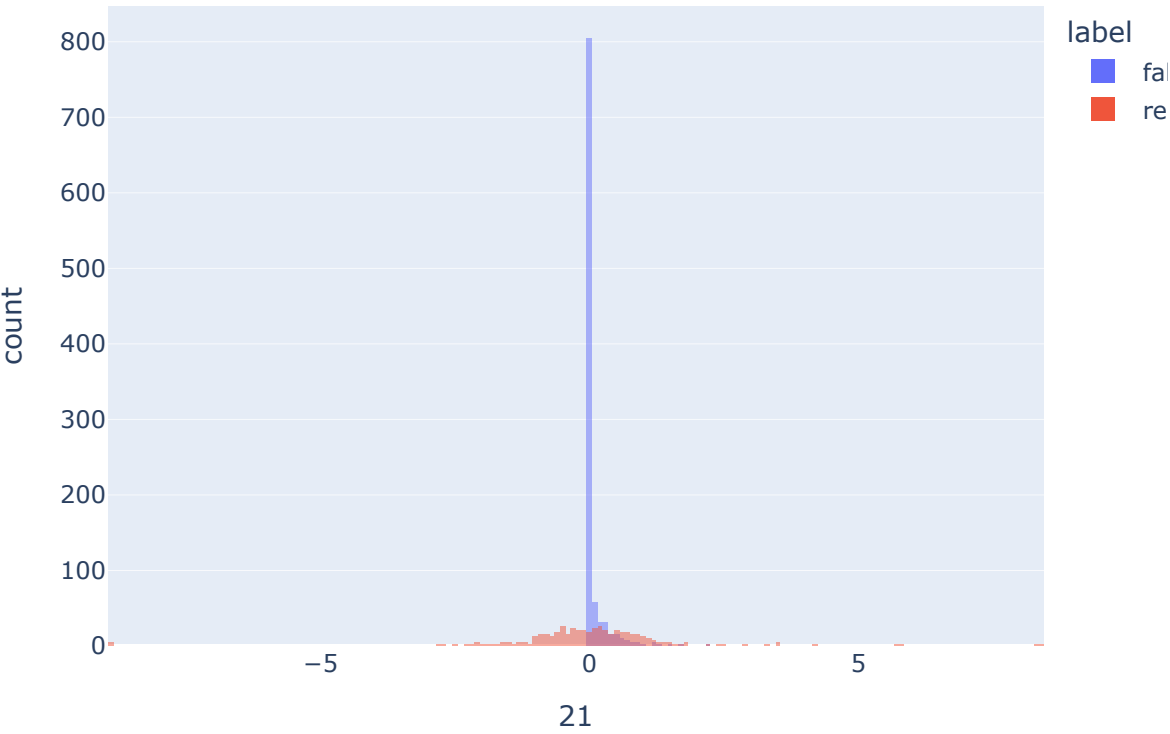
```
/var/folders/w5/_kz_512d0bd76w88dw4ym66r0000gn/T/ipykernel_14370/3618524353.py:2: RuntimeWarning:
```

More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`). Consider using `matplotlib.pyplot.close`.

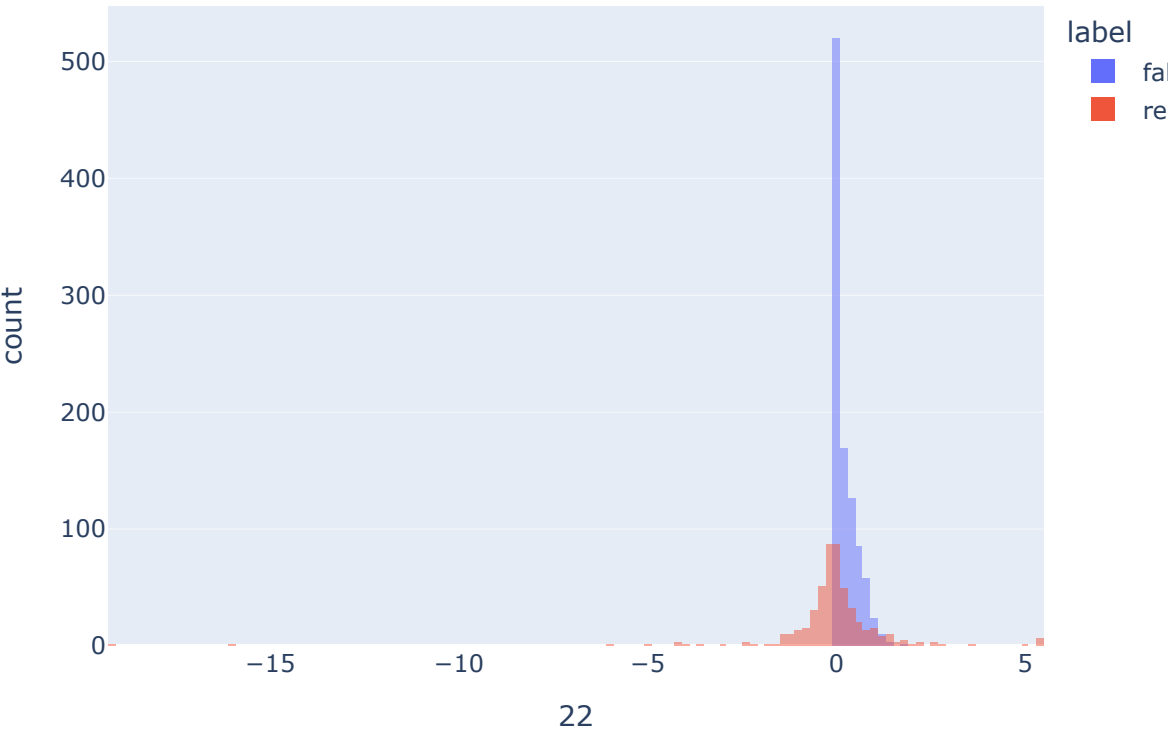
Feature 20



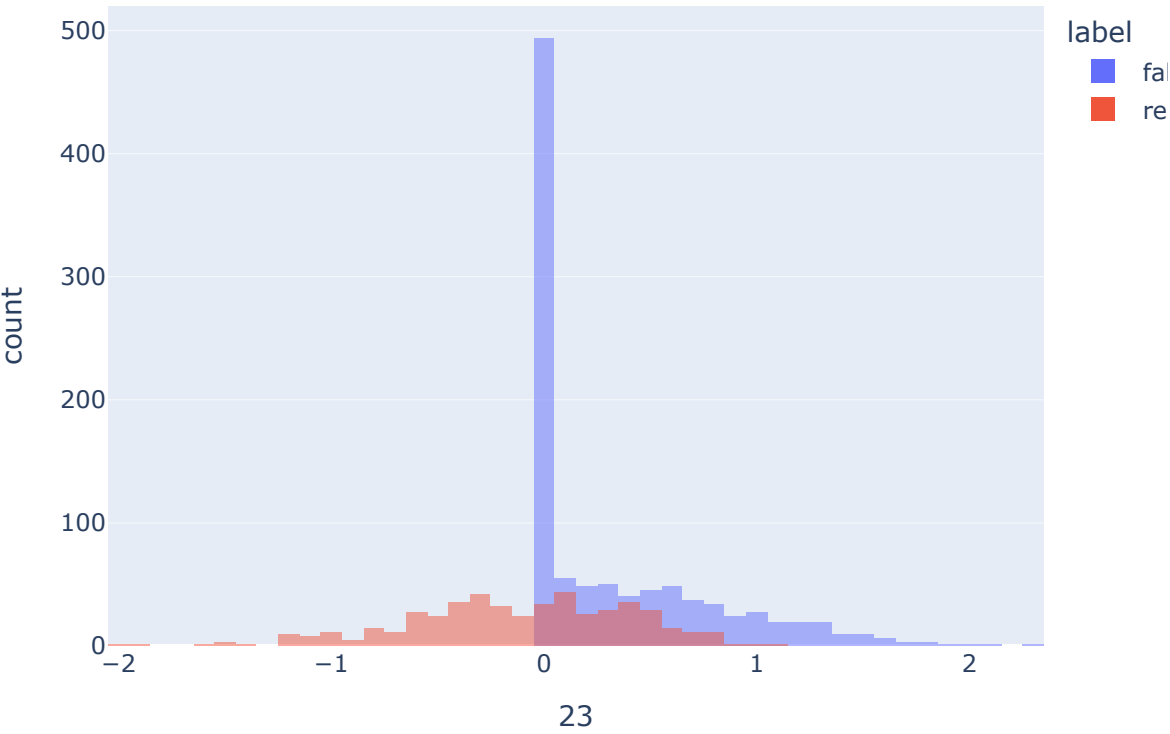
Feature 21



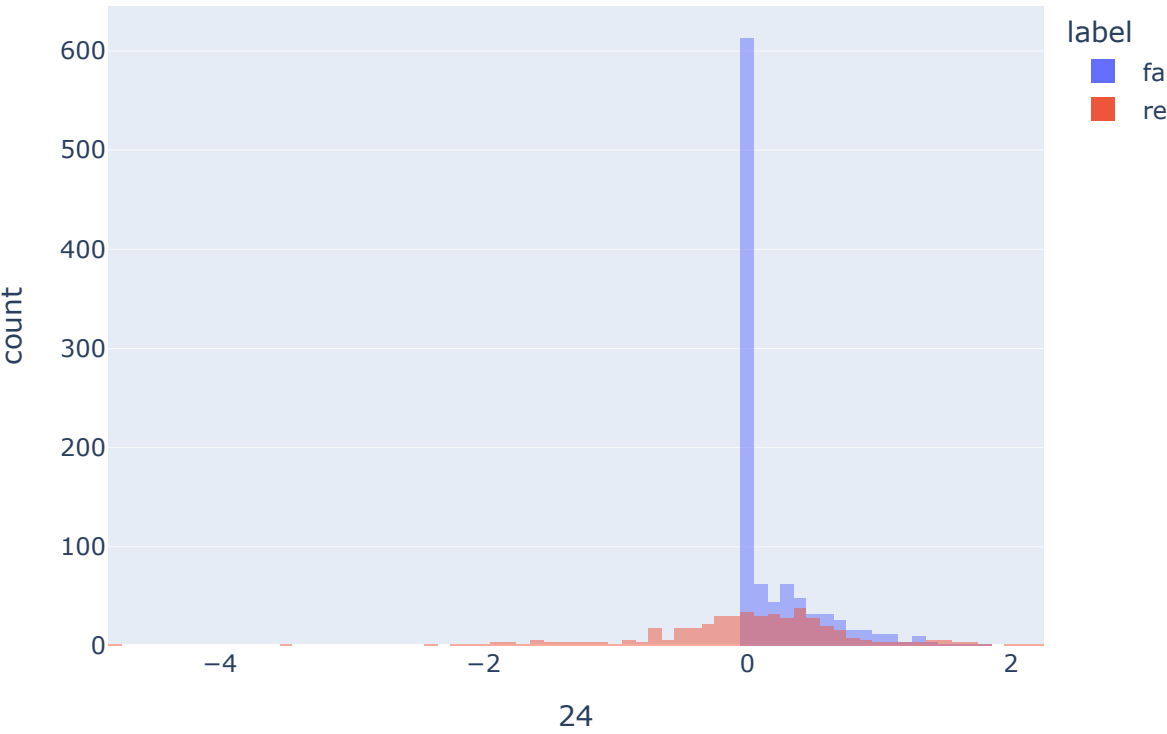
Feature 22



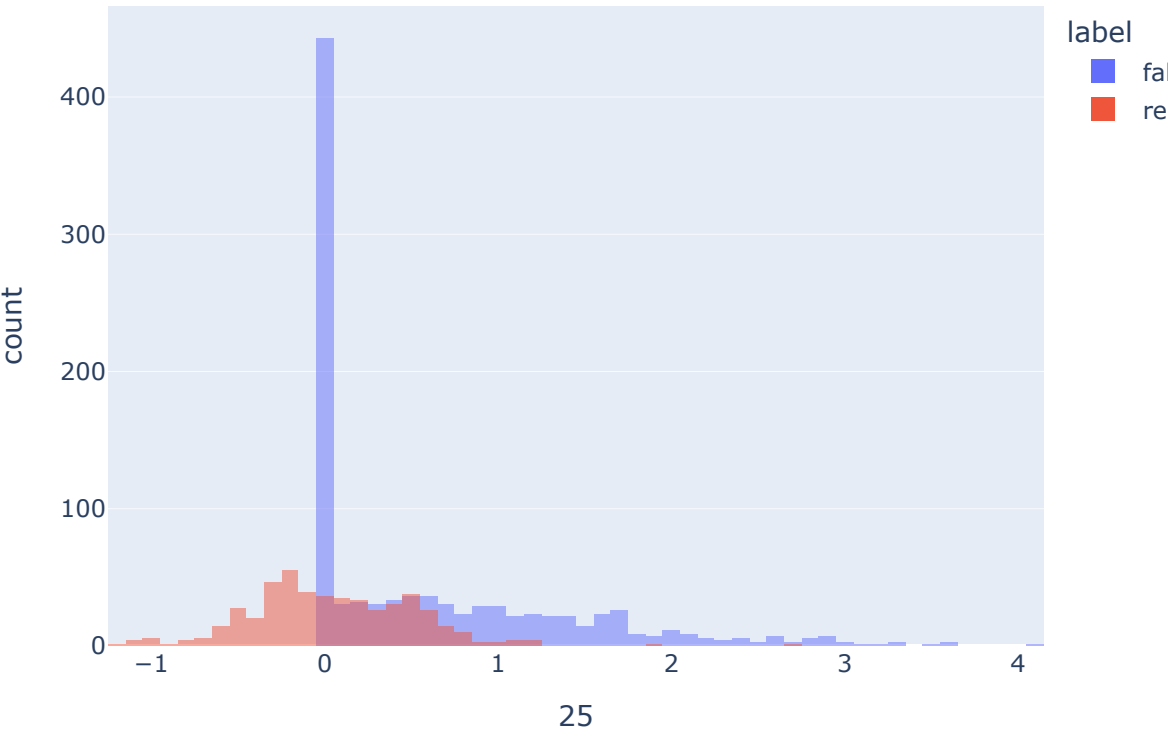
Feature 23



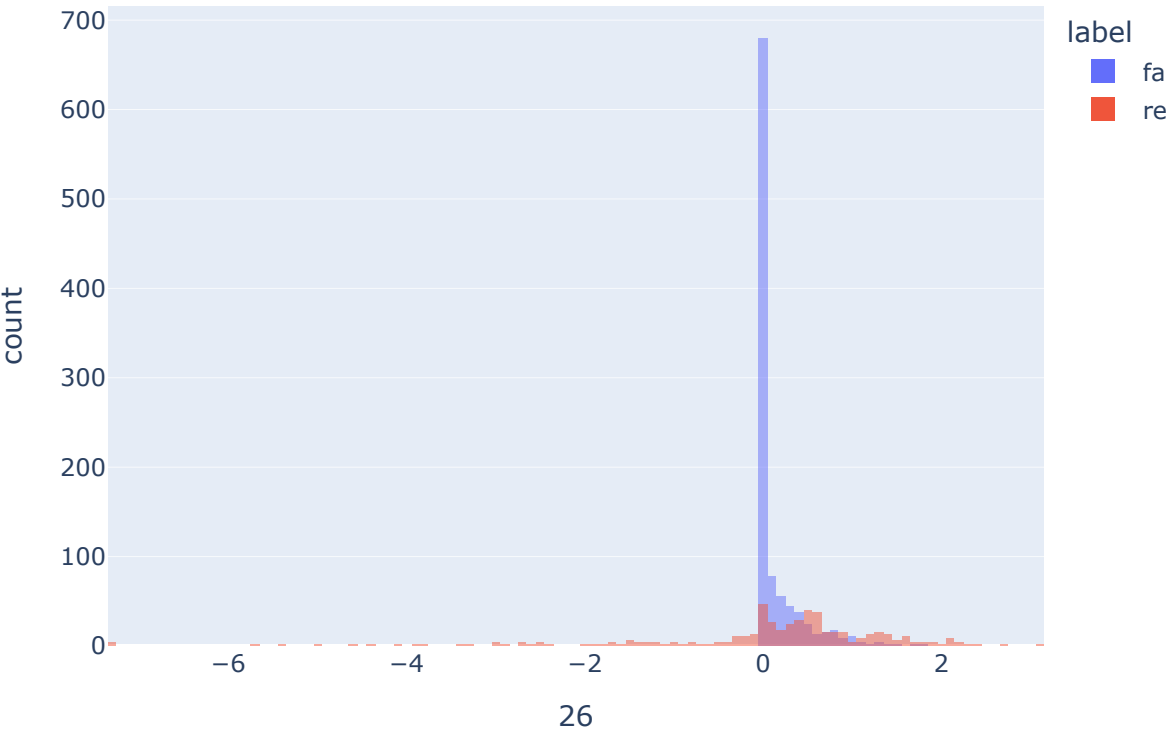
Feature 24



Feature 25

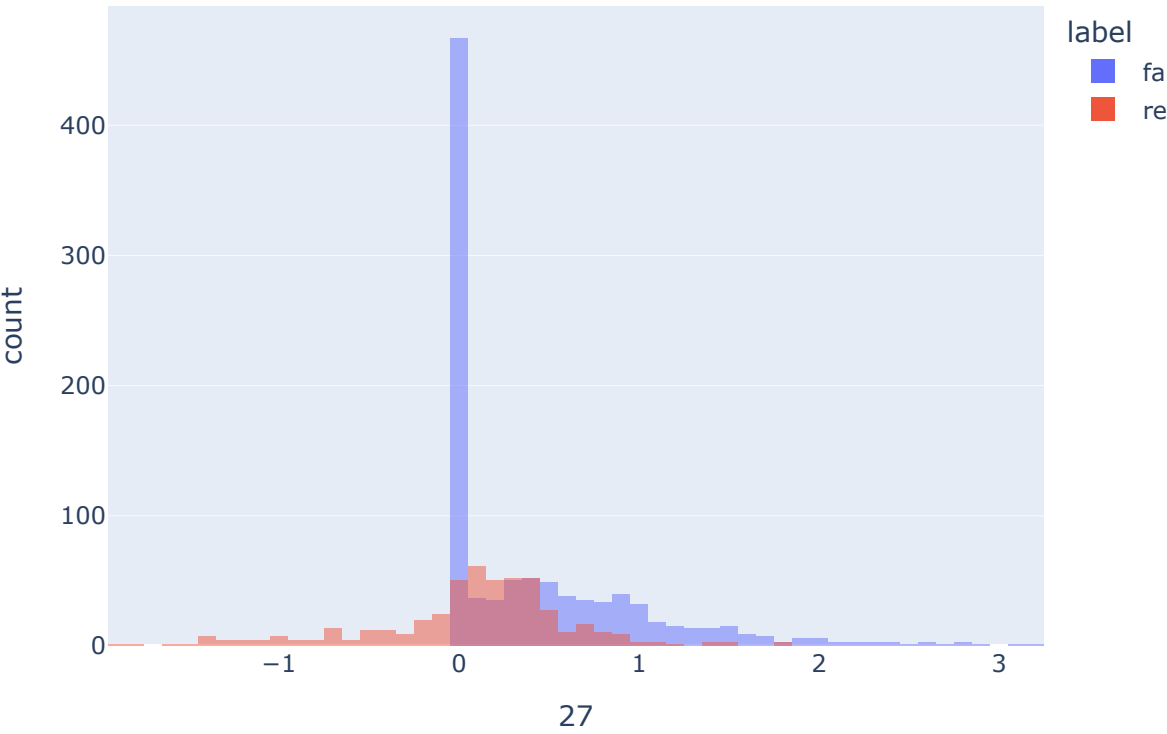


Feature 26

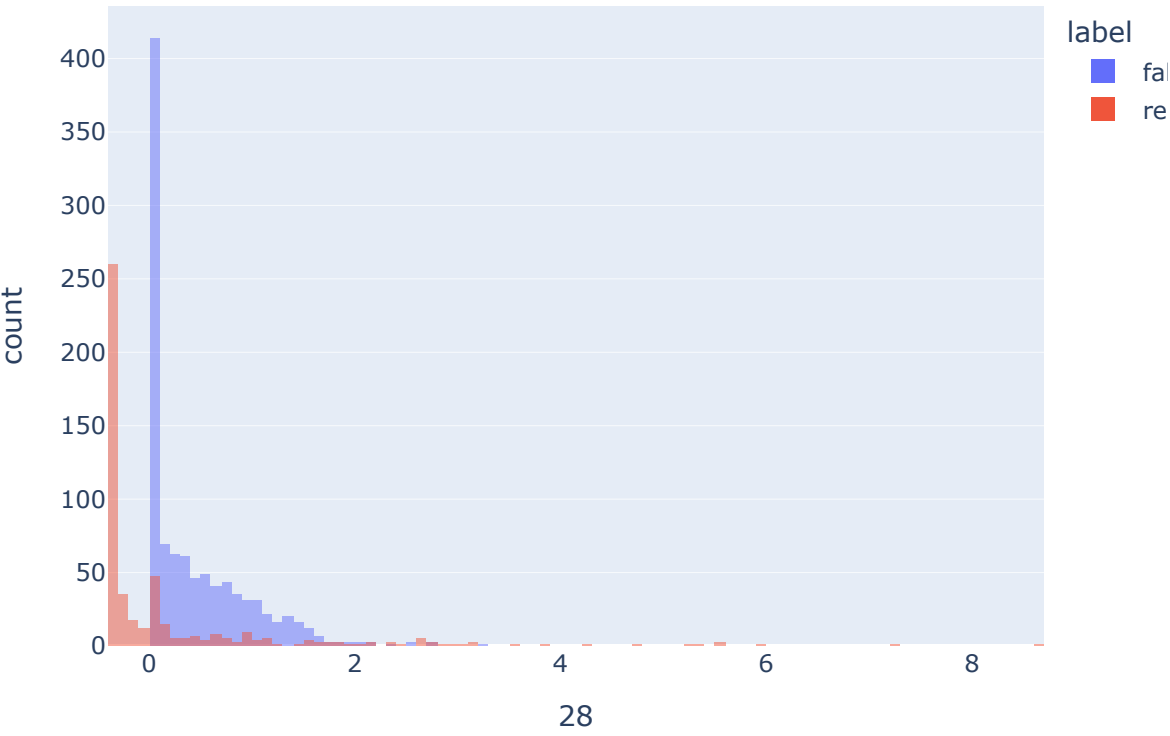




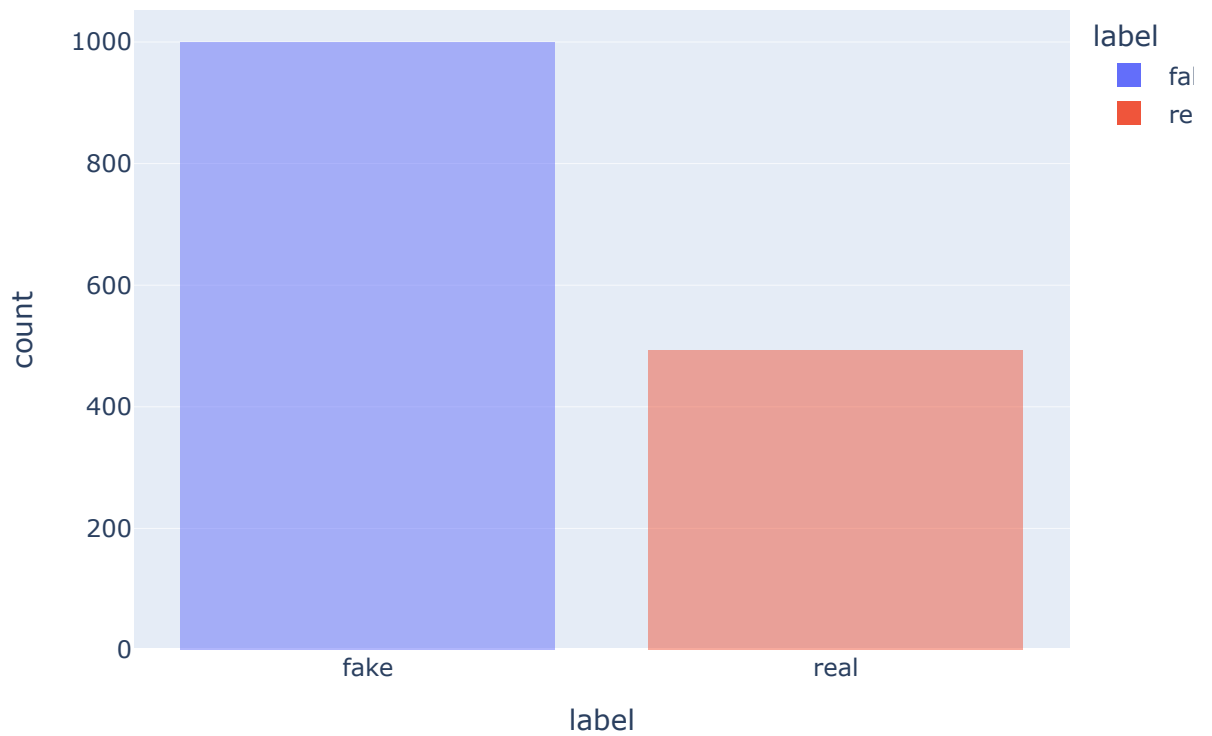
Feature 27



Feature 28



## Feature label

[illegible]

<Figure size 640x480 with 0 Axes>  
<Figure size 640x480 with 0 Axes>  
<Figure size 640x480 with 0 Axes>  
<Figure size 640x480 with 0 Axes>

In [ ]: